

CS405 COMPUTER SYSTEM ARCHITECTURE

By Elizabeth Isaac

**Department of Computer Science and
engineering**

M A College of Engineering


**K. Hwang and Naresh Jotwani, Advanced Computer Architecture,
Parallelism, Scalability, Programmability, TMH, 2010**

Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - Architectural development tracks
 - Conditions of parallelism



ABACUS




ROTATING WHEEL CALCULATOR



DIFFERENCE ENGINE







COMING UP TECHNOLOGY

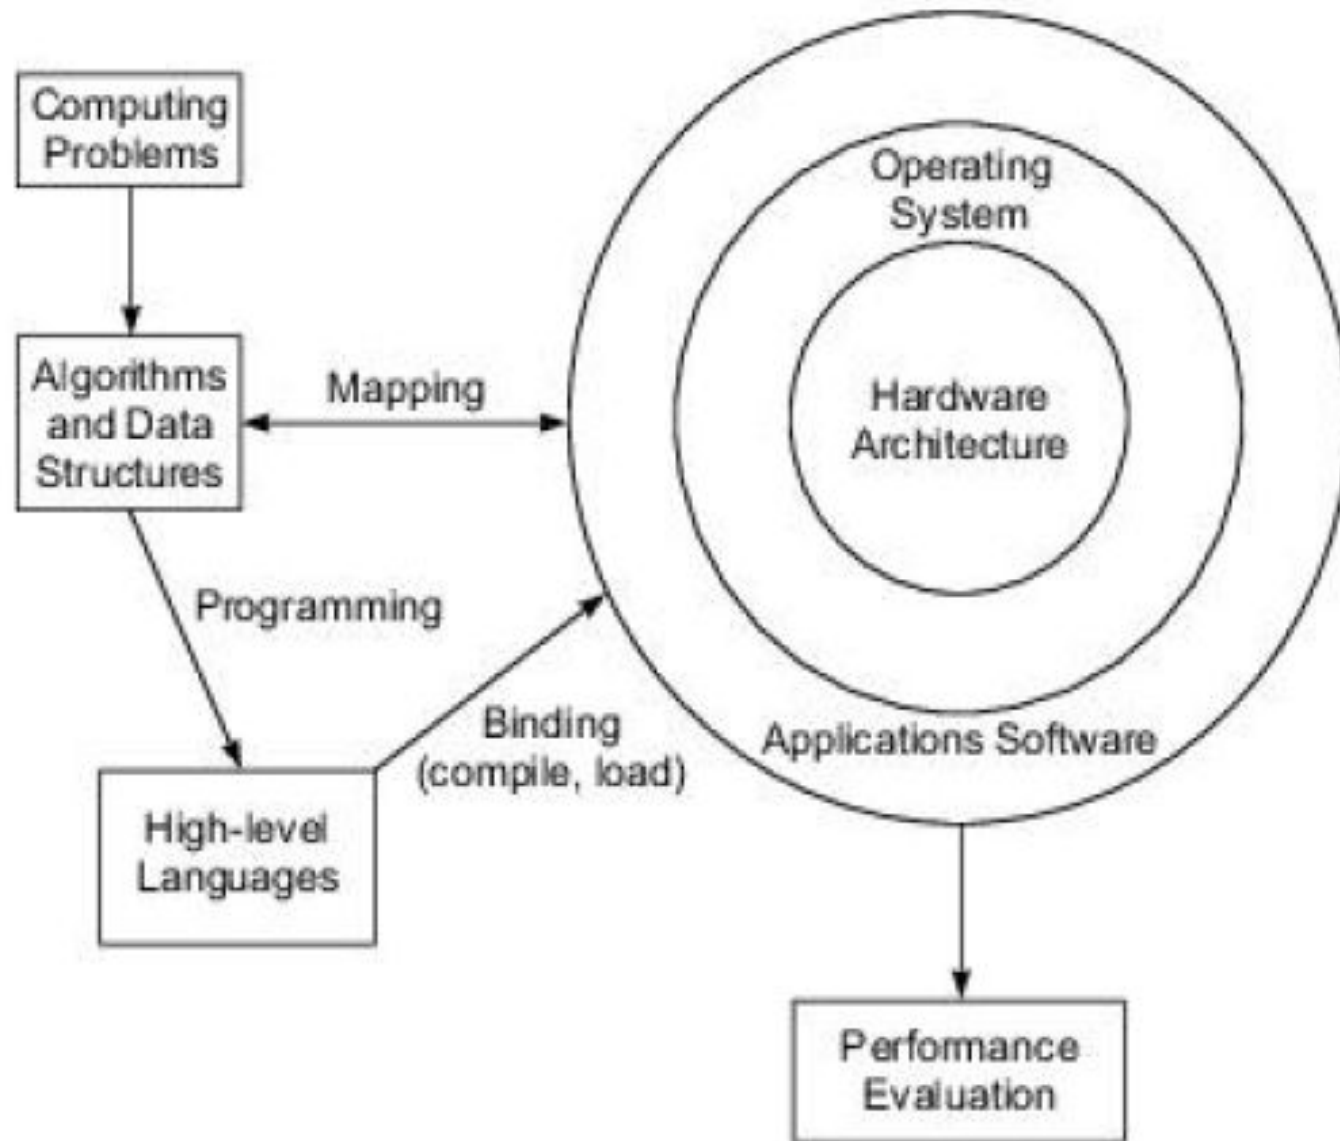


Parallel computer models

Computer Generations : Electronic computers have gone through five generations of development.

<i>Generation</i>	<i>Technology and Architecture</i>	<i>Software and Applications</i>	<i>Representative Systems</i>
First (1945–54)	Vacuum tubes and relay memories, CPU driven by PC and accumulator, fixed-point arithmetic.	Machine/assembly languages, single user, no subroutine linkage, programmed I/O using CPU.	ENIAC, Princeton IAS, IBM 701.
Second (1955–64)	Discrete transistors and core memories, floating-point arithmetic, I/O processors, multiplexed memory access.	HLL used with compilers, subroutine libraries, batch processing monitor.	IBM 7090, CDC 1604, Univac LARC.
Third (1965–74)	Integrated circuits (SSI/-MSI), microprogramming, pipelining, cache, and lookahead processors.	Multiprogramming and time-sharing OS, multiuser applications.	IBM 360/370, CDC 6600, TI-ASC, PDP-8.
Fourth (1975–90)	LSI/VLSI and semiconductor memory, multiprocessors, vector supercomputers, multicomputers.	Multiprocessor OS, languages, compilers, and environments for parallel processing.	VAX 9000, Cray X-MP, IBM 3090, BBN TC2000.
Fifth (1991–present)	Advanced VLSI processors, memory, and switches, high-density packaging, scalable architectures.	Superscalar processors, systems on a chip, massively parallel processing, grand challenge applications, heterogeneous processing.	See Tables 1.3–1.6 and Chapter 13.

- **Elements of Modern Computers**



- **Computing Problem**

- computer architecture concept is **no longer restricted** to the bare machine hardware structure
- Consist of machine hardware, an instruction set, system software, application programs, and user interfaces
- **Cost effective solutions**
- Example:
 - (1) **Alpha numerical problems in business and government** (solutions demand efficient transaction processing, large database management and information retrieval operations)
 - (2) **Artificial intelligence [AI] problems** (solution demand logic inferences and symbolic manipulations)

- **Algorithm and Data Structures**

- Special algorithms (deterministic and non-deterministic) and data structures
- Problem formulation and the development of parallel algorithms

- **Hardware Resources**

- Three nested circles in figure
- Special hardware interfaces
- Software interface programs

- **Operating Systems**

- Allocation and Deallocation
- Performance Evaluation
- Mapping and Optimal mappings

- **System Software Support**

- Efficient programs in high-level languages



- **Compiler Support**

3 compiler upgrades:

- **Preprocessor, Precompiler and Parallelizing Compiler**
- **Preprocessor** : Uses a sequential compiler and a low-level library of the target computer
- **Precompiler** : Requires Program flow analysis, dependence checking and limited optimizations
- **Parallelizing Compiler** : Automatically detect parallelism in source code

Module 1

- **Parallel computer models**
 - **Evolution of Computer Architecture**
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - Architectural development tracks
 - Conditions of parallelism

Von Neumann Architecture

**Von Neumann
Architecture?**

**Input
Device**

Central Processing Unit

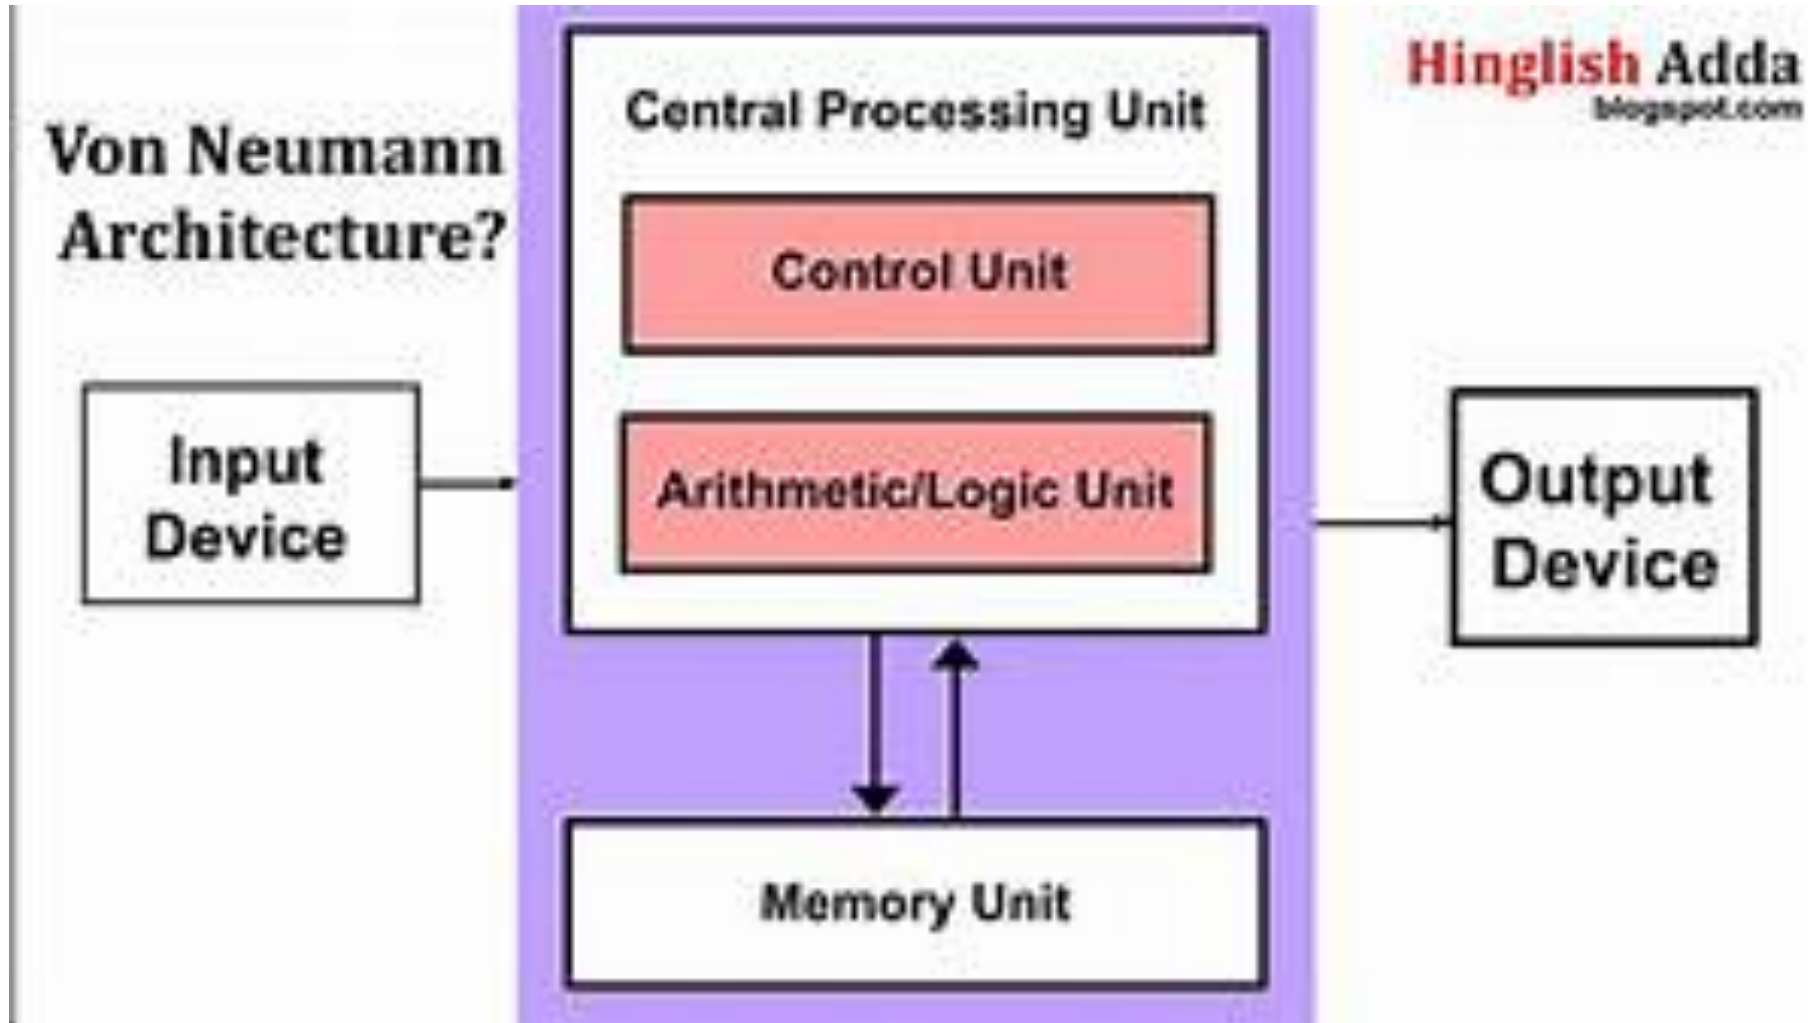
Control Unit

Arithmetic/Logic Unit

Memory Unit

**Output
Device**

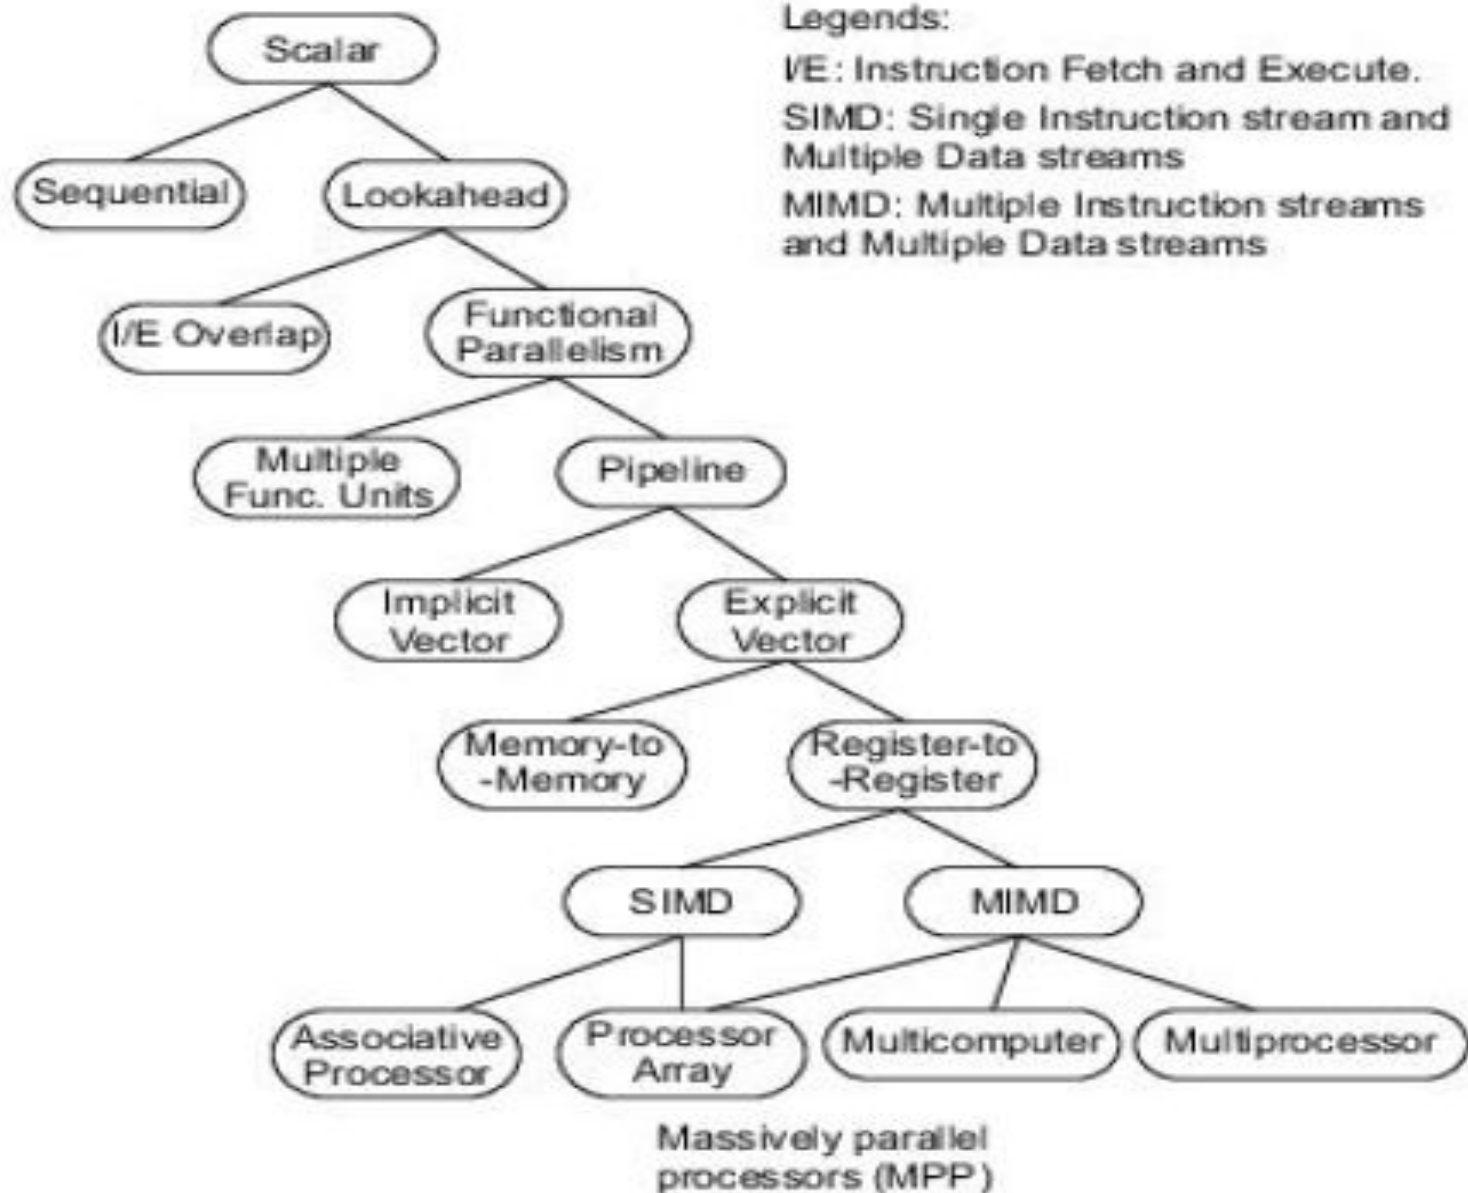
Hinglish Adda
blogspot.com



Evolution of Computer Architecture


- Hardware organization and programming/software requirements
- Abstracted by instruction set
- Architectural evolution from sequential scalar computers to vector processors and parallel computers

• Architectural Evolution




- **Flynn's Classification**

- Michael Flynn (1972)
- Based on notations of instructions and data streams
 - SISD (Single instruction streams over single data streams)
 - SIMD (Single instruction streams over multiple data streams)
 - MISD (Multiple instruction streams over single data streams)
 - MIMD (Multiple instruction streams over multiple data streams)



(a) SISD uniprocessor architecture



(b) SIMD architecture (with distributed memory)

Captions:

CU = Control Unit

PU = Processing Unit


MU = Memory Unit

IS = Instruction Stream

DS = Data Stream

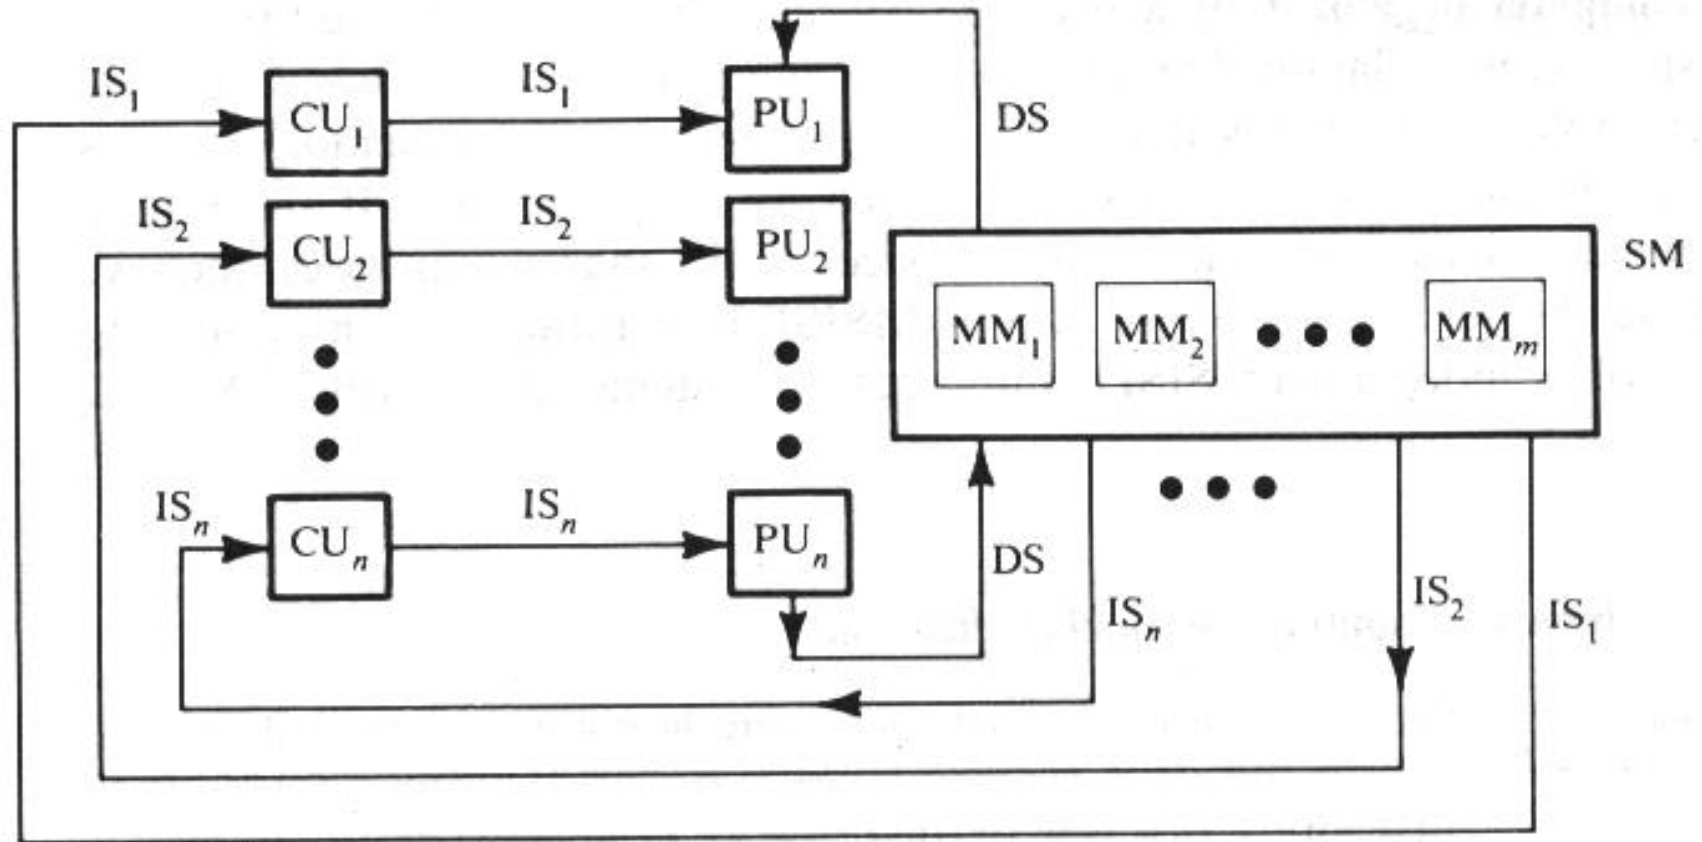
PE = Processing Element

LM = Local Memory

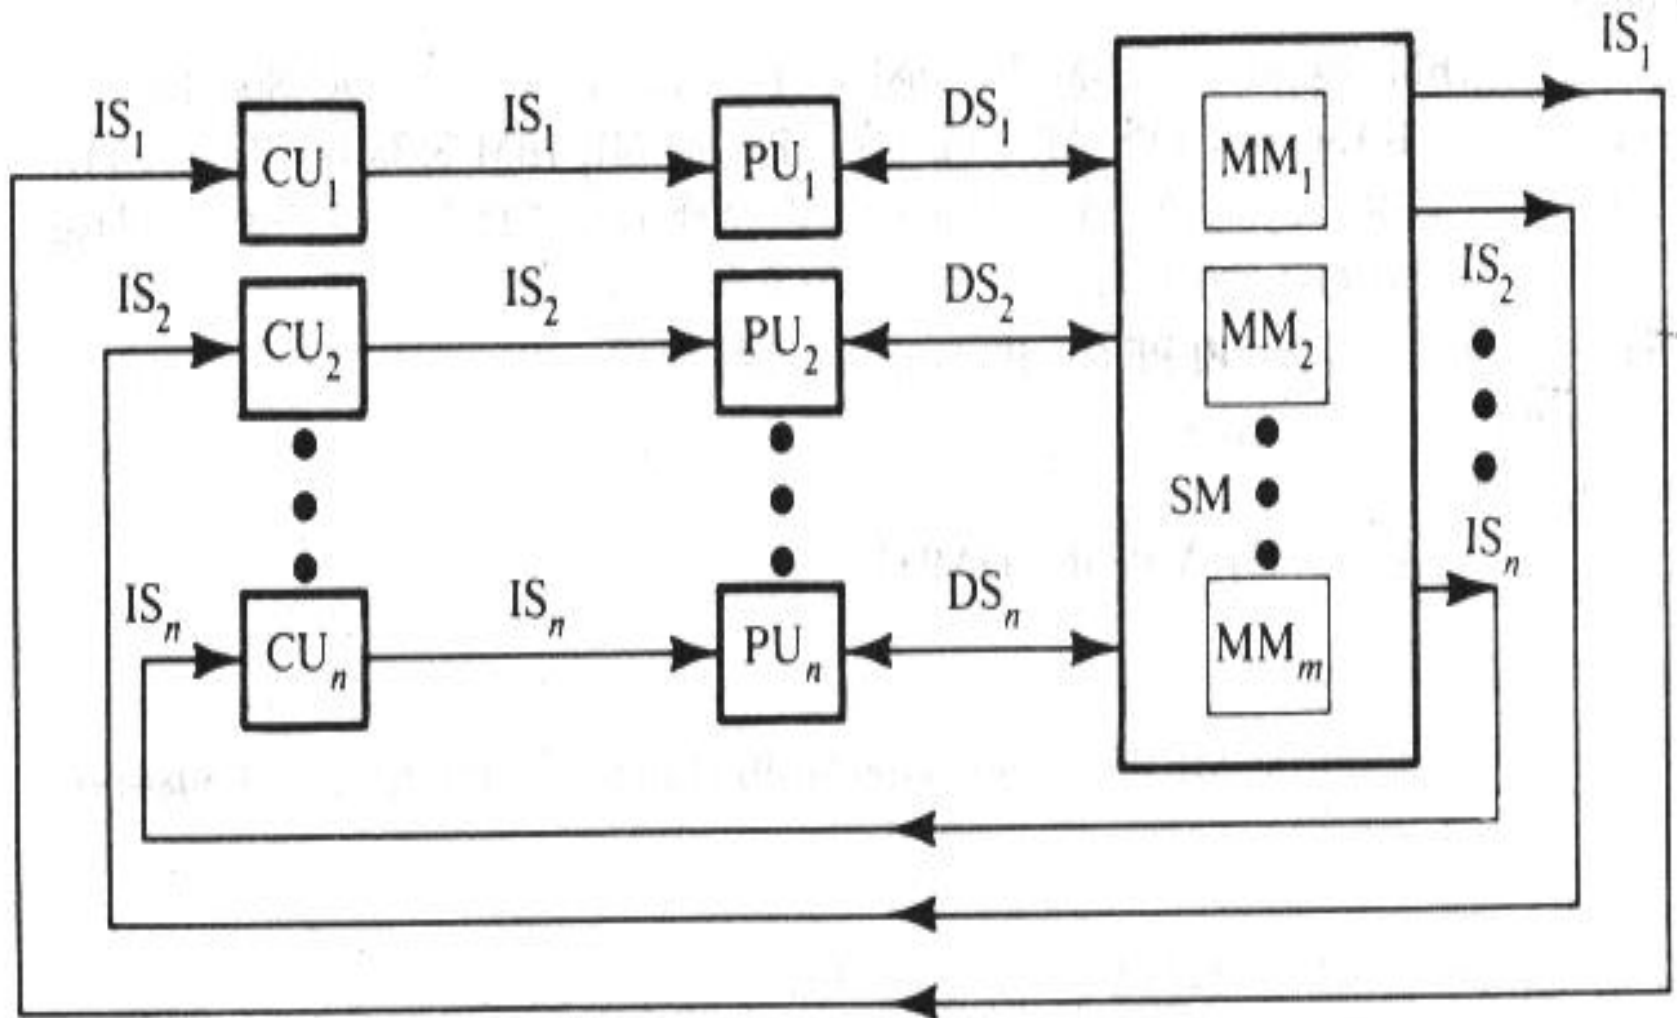


(c) MIMD architecture (with shared memory)

MISD




MIMD



- Parallel/ vector Computer:
 - Programs in MIMD mode
 - Two classes
 - Shared-Memory Multiprocessors
 - Message- passing multicomputers

- Six layers for computer system development



Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - **System Attributes to performance**
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - Architectural development tracks
 - Conditions of parallelism

System Attributes to Performance

- Machine capability and program behavior
- **Clock Rate and CPI:**
 - Processor is driven by a clock with a constant cycle time T . The inverse of the cycle time is the clock rate ($f = 1/T$).
 - Size by program Instruction Count(I_c)
 - Cycles per instruction (CPI)

System Attributes to Performance

- **Performance Factors**

I_c : number of instructions(Instruction count)

CPU Time, $T = I_c * CPI * T$ (T in Seconds/Program)

- we can rewrite as $T = I_c * (p + m * k) * T$

p : Number of processor cycles for ID and EX

m : Number of memory references needed

K : Ratio between memory cycle and processor cycle

System Attributes to Performance

- **System Attribute:** five performance factors (I_c , p , m , k , T) are influenced by four system attributes:
 - Instruction-set architecture
 - Compiler technology
 - CPU implementation and control
 - Cache and memory hierarchy

- **Performance Factors Versus System Attributes**

System Attributes	Performance Factors				
	Instr. Count, I_c	Average Cycles per Instruction, CPI			Processor Cycle Time, τ
		Processor Cycles per Instruction, p	Memory References per Instruction, m	Memory-Access Latency, k	
Instruction-set Architecture	✓	✓			
Compiler Technology	✓	✓	✓		
Processor Implementation and Control		✓			✓
Cache and Memory Hierarchy				✓	✓

System Attributes to Performance

- MIPS Rate

$$\text{MIPS rate} = \frac{Ic}{T * 10^6}$$

$$T = \frac{Ic * CPI}{f}$$

$$\text{MIPS rate} = \frac{Ic * f}{Ic * CPI * 10^6}$$

$$\text{MIPS rate} = \frac{f}{CPI * 10^6}$$

$$CPI = \frac{C}{Ic}$$

$$\text{MIPS rate} = \frac{f * Ic}{C * 10^6}$$

THROUGHPUT rate

$$W_p = \frac{f}{I_c * CPI} = \frac{MIPS * 10^6}{I_c}$$

- **Problem:**

- Machine A runs a program in 20 seconds. Machine B runs the same program in 25 seconds. How many times faster is machine A?

$$n = \text{ExecTime}(B) / \text{ExecTime}(A)$$

$$= \frac{25 \text{ sec}}{20 \text{ sec}} = 1.25$$

Machine A is 1.25 times faster than Machine B

Numerical:- A benchmark program is executed on a 40MHz processor. The benchmark program has the following statistics.

Instruction Type	Instruction Count	Clock Cycle Count
Arithmetic	45000	1
Branch	32000	2
Load/Store	15000	2
Floating Point	8000	2

Calculate for the above program
average CPI, MIPS rate & execution time

$$\text{Average CPI} = \frac{C}{I_c}$$

$$\frac{C}{I_c} = \frac{\text{Total \# cycles to execute a whole program}}{\text{Total Instructions}}$$

$$= \frac{45000 \times 1 + 32000 \times 2 + 1500 \times 2 + 8000 \times 2}{45000 + 3200 + 15000 + 8000}$$

$$= \frac{155000}{100000}$$

$$\text{CPI} = \underline{\underline{1.55}}$$

$$\underline{\underline{\text{Execution Time} = C / f}}$$

$$T = 155000 / 40 \times 10^6$$

$$T = 0.155 / 40$$

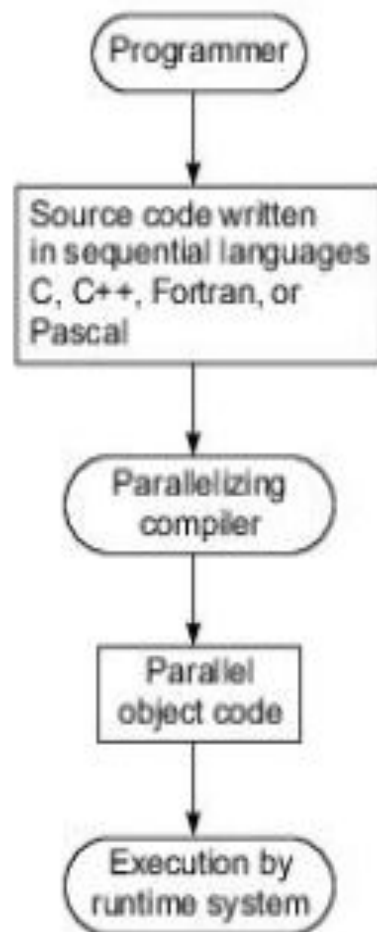
$$T = \underline{\underline{3.875 \text{ ms}}}$$

$$\text{MIPS rate} = I_c / T \times 10^6$$

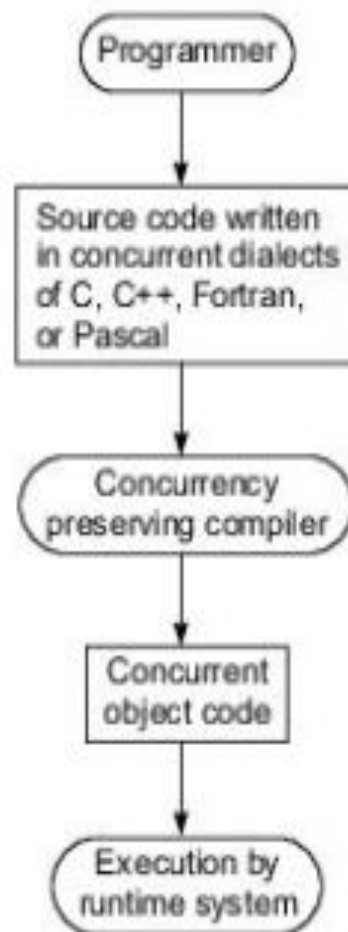
$$= 100000 / 40 \times 10^6$$

$$\text{MIPS rate} = \underline{\underline{25.81}}$$

- **Implicit parallelism** : a conventional language, such as C, C++, fortran, or pascal to write the source program
 - **Explicit parallelism** : requires more effort by the programmer to develop a source program using parallel dialects of C, C++, fortran, or pascal
 - Parallelism is explicitly specified in the user programs
 - The compiler needs to preserve parallelism and, where possible, assigns target machine resources.
- New programming language chapel



(a) Implicit parallelism



(b) Explicit parallelism

Two approaches to parallel programming (Courtesy of Charles Seitz; adapted with permission from "Concurrent Architectures", p. 51 and p. 53, *VLSI and Parallel Computation*, edited by Suaya and Birtwistle, Morgan Kaufmann Publishers, 1990)

Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - **Amdahl's law for a fixed workload**
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - Architectural development tracks
 - Conditions of parallelism


Amdahl's Law for a Fixed Workload

- Number of processors increases in a parallel computer, the fixed load is distributed to more processors for parallel execution
- Minimal turnaround time is the primary goal.
- Speedup obtained for time-critical applications is called fixed—load speedup

Amdahl's Law

- A program (or algorithm) which can be parallelized can be split up into two parts:
 - A part which cannot be parallelized and
 - A part which can be parallelized


First of all, a program can be broken up into a **non-parallelizable part B**, and a **parallelizable part 1-B**, as illustrated by this diagram: **total execution time T(1)**.



Execution time with a **parallelization factor of 2:**



Execution time with a **parallelization factor of 3:**



Amdahl's law

Amdahl's Law (1967)

For a given problem size, the speedup does not increase linearly as the number of processors increases. In fact, the speedup tends to become saturated.

This is a consequence of Amdahl's Law.

According to Amdahl's Law, a program contains two types of operations:

Completely sequential

Completely parallel

Let, the time **T_s** taken to perform sequential operations be a fraction **α** ($0 < \alpha \leq 1$) of the total execution time **T(1)** of the program, then the time **T_p** to perform parallel operations shall be **(1- α)** of **T(1)**

Amdahl's Law

Thus, $T_s = \alpha.T(1)$ and $T_p = (1-\alpha).T(1)$

Assuming that the parallel operations achieve linear speedup (i.e. these operations use $1/n$ of the time taken to perform on each processor), then

$$T(n) = T_s + T_p/n = \alpha.T(1) + \frac{(1-\alpha).T(1)}{n}$$

Thus, the speedup with n processors will be:

$$S(n) = \frac{T(1)}{T(n)} = \frac{T(1)}{\alpha.T(1) + \frac{(1-\alpha).T(1)}{n}} = \frac{1}{\alpha + \frac{(1-\alpha)}{n}}$$

$$S_n = \frac{n}{1 + (n-1)\alpha}$$

- **Example: 1**

- Suppose that a calculation has a 4% serial portion,
- a) What is the limit of speedup on 16 processors?
- b) What is the maximum speedup?

Ans:

a) Limit of speedup on 16 processors
= $16 / (1 + (16 - 1) * .04) = 10$

b) The maximum speedup = $1/\alpha$
= $1/0.04 = 25$

$$S_n = \frac{n}{1 + (n - 1)\alpha}$$

Amdahl's Law

- The performance gain that can be obtained by improving some portion of a computer can be calculated using Amdahl's law.
- Amdahl's law states that the performance improvement to be gained from using some faster mode of execution is limited by the fraction of the time the faster mode can be used.
- Law defines the term 'speedup'.

$$\text{Speedup} = \frac{\text{Performance for entire task using the enhancement when possible}}{\text{Performance for entire task without using the enhancement}}$$

Alternatively,

$$\text{Speedup} = \frac{\text{Execution time for entire task without using the enhancement}}{\text{Execution time for entire task using the enhancement when possible}}$$

Amdahl's Law

- The execution time using the original computer with the enhanced mode will be the time spent using the unenhanced portion of the computer

$$\text{Execution time}_{\text{new}} = \text{Execution time}_{\text{old}} \times \left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)$$

The overall speedup is the ratio of the execution times:

$$\text{Speedup}_{\text{overall}} = \frac{\text{Execution time}_{\text{old}}}{\text{Execution time}_{\text{new}}} = \frac{1}{(1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}}}$$

3. We are considering an enhancement to the processor of a web server. The new CPU is 20 times faster on search queries than the old processor. The old processor is busy with search queries 70% of the time, what is the speedup gained by integrating the enhanced CPU?

$$Speedup = \frac{1}{(1 - Fraction_{enhanced}) + \frac{Fraction_{enhanced}}{Speedup_{enhanced}}}$$

$$Fraction_{enhanced} = 70\% = 0.70$$

$$Speedup_{enhanced} = 20$$

$$Speedup = \frac{1}{(1 - 0.70) + \frac{0.70}{20}} = \frac{1}{0.335} = 2.985$$

Parallel programs

- If 30% of the execution time may be the subject of a speedup, p will be 0.3; if the improvement makes the affected part twice as fast, s will be 2. Amdahl's law states that the overall speedup of applying the improvement will be:

$$S_{\text{latency}} = \frac{1}{1 - p + \frac{p}{s}} = \frac{1}{1 - 0.3 + \frac{0.3}{2}} = 1.18.$$

- For example, assume that we are given a serial task which is split into four consecutive parts, whose percentages of execution time are $p1 = 0.11$, $p2 = 0.18$, $p3 = 0.23$, and $p4 = 0.48$ respectively. The 1st part is not speed up, so $s1 = 1$, while the 2nd part is sped up 5 times, so $s2 = 5$, the 3rd part is sped up 20 times, so $s3 = 20$, and the 4th part is sped up 1.6 times, so $s4 = 1.6$.

- By using Amdahl's law, the overall speedup is

$$S_{\text{latency}} = \frac{1}{\frac{p_1}{s_1} + \frac{p_2}{s_2} + \frac{p_3}{s_3} + \frac{p_4}{s_4}} = \frac{1}{\frac{0.11}{1} + \frac{0.18}{5} + \frac{0.23}{20} + \frac{0.48}{1.6}} = 2.19.$$

- Notice how the 20 times and 5 times speedup on the 2nd and 3rd parts respectively don't have much effect on the overall speedup when the 4th part (48% of the execution time) is accelerated by only 1.6 times.

Module 1

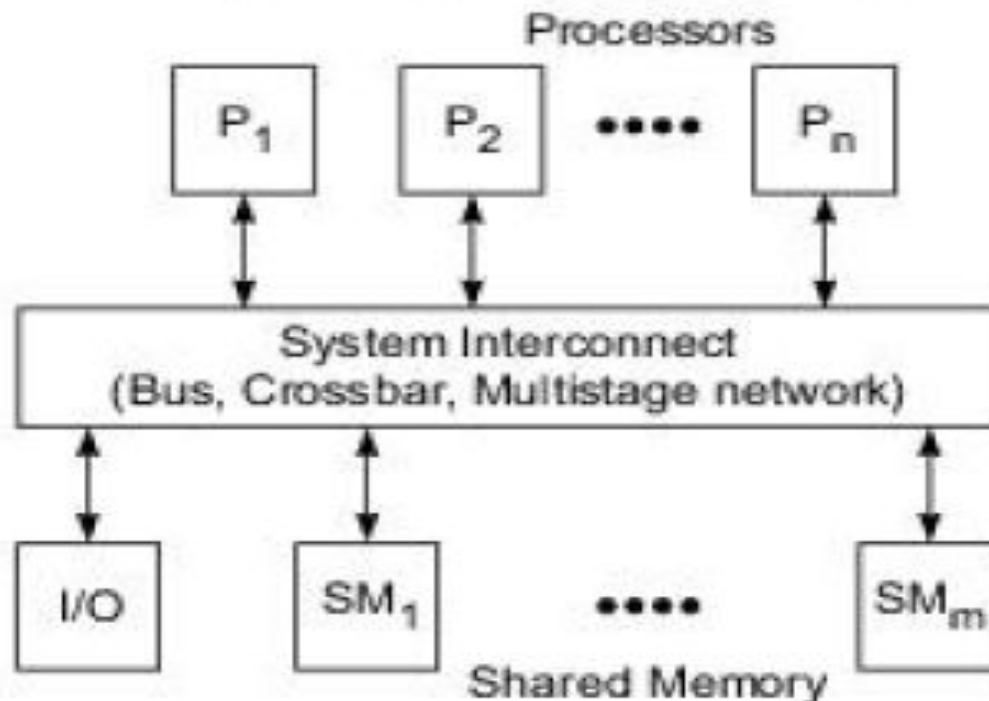
- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - **Multiprocessors and Multicomputers**
 - Multivector and SIMD computers
 - Architectural development tracks
 - Conditions of parallelism

Multiprocessors and Multicomputers

- 3 shared- memory multiprocessor models
 - Uniform memory-access(UMA) Model
 - Non-uniform memory-access(NUMA) Model
 - Cache-only memory architecture(COMA) Model

- **UMA Model**


- Physical memory is uniformly shared by all the processors
- Equal access time – so called uniform memory access
- Private cache



Multiprocessors – Tightly coupled systems

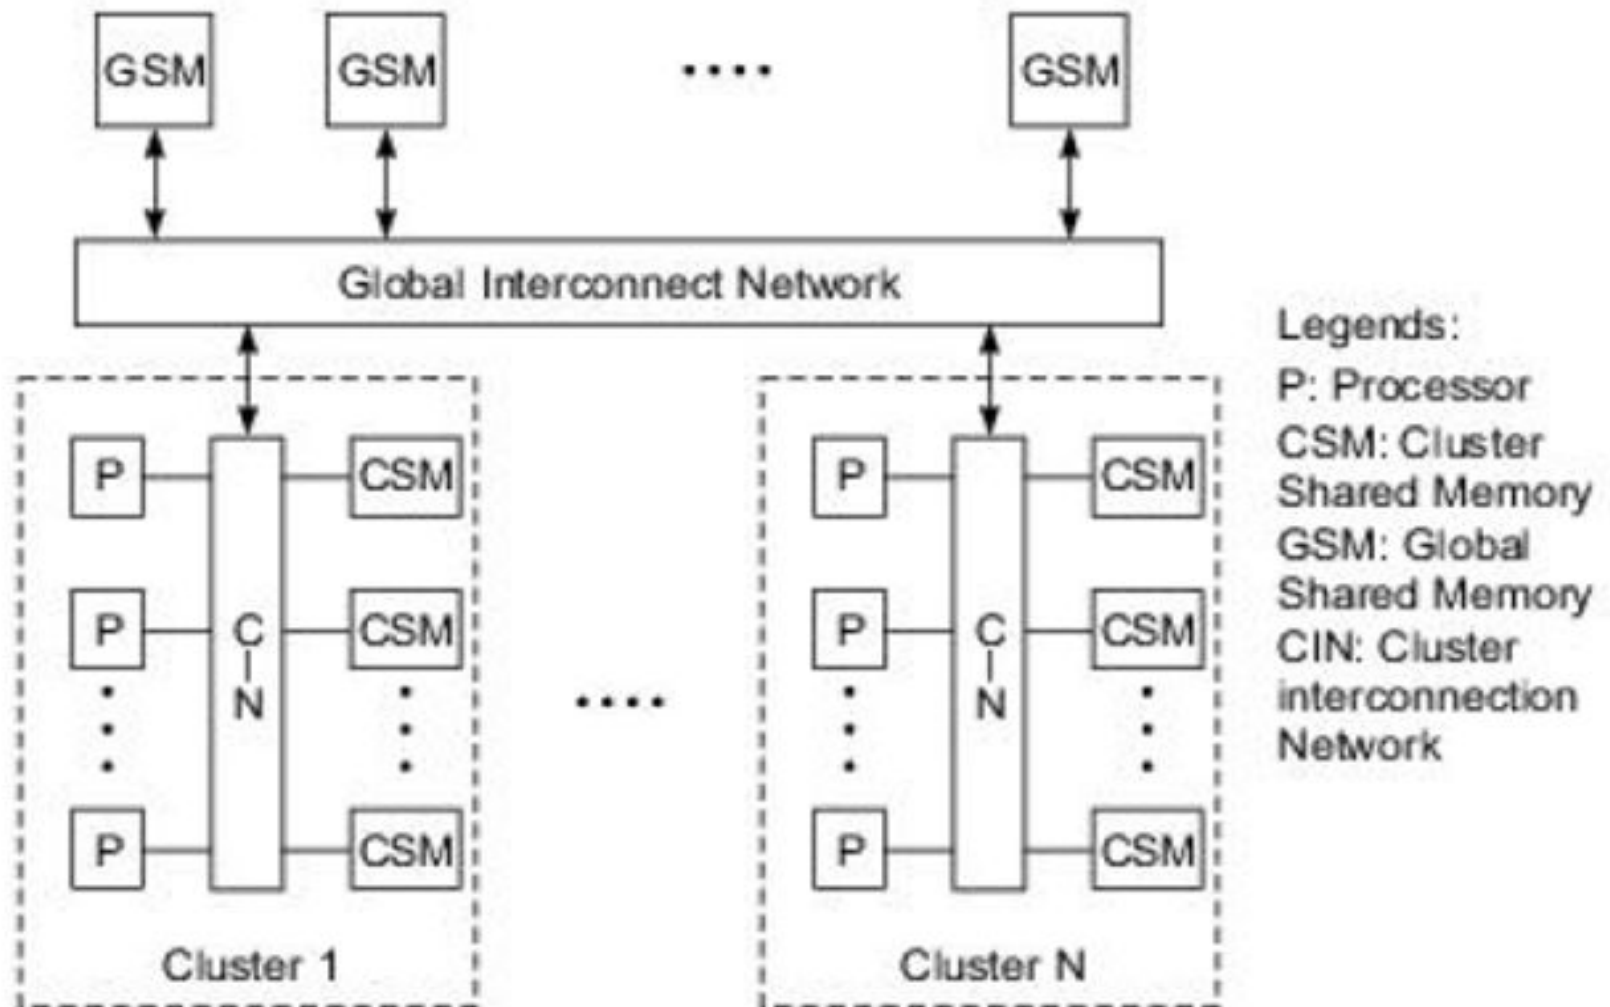
- General-purpose and time sharing applications
- Parallel events, synchronization and communication using shared variables
- Symmetric multiprocessor
- Asymmetric multiprocessor

- **NUMA Model**
 - Shared-memory system



(a) Shared local memories (e.g. the N Butterfly)


NUMA Model



(b) A hierarchical cluster model (e.g. the Cedar system at the University of Illinois)

- **COMA Model**

- Special case of a NUMA machine
- Distributed main memories are converted to caches
- No memory hierarchy
- Caches form a global address space
- Data will eventually migrate to where it will be used.



The COMA model of a multiprocessor (P: Processor, C: Cache, D: Directory; e.g. the KSR-1)

Other Variations for multiprocessor

- Cache-coherent non-uniform memory access model (CC-NUMA)


- Representative Multiprocessors

<i>Company and Model</i>	<i>Hardware and Architecture</i>	<i>Software and Applications</i>	<i>Remarks</i>
Sequent Symmetry S-81	Bus-connected with 30 i386 processors, IPC via SLIC bus; Weitek floating-point accelerator.	DYNIX/OS, KAP/Sequent preprocessor, transaction multiprocessing.	Latter models designed with faster processors of the family.
IBM ES/9000 Model 900/VF	6 ES/9000 processors with vector facilities, crossbar connected to I/O channels and shared memory.	OS support: MVS, VM KMS, AIX/370, parallel Fortran, VSF V2.5 compiler.	Fiber optic channels, integrated cryptographic architecture.
BBN TC-2000	512 M88100 processors with local memory connected by a Butterfly switch, a NUMA machine.	Ported Mach/OS with multiclustering, parallel Fortran, time-critical applications.	Latter models designed with faster processors of the family.

- **Distributed-Memory Multicomputers**

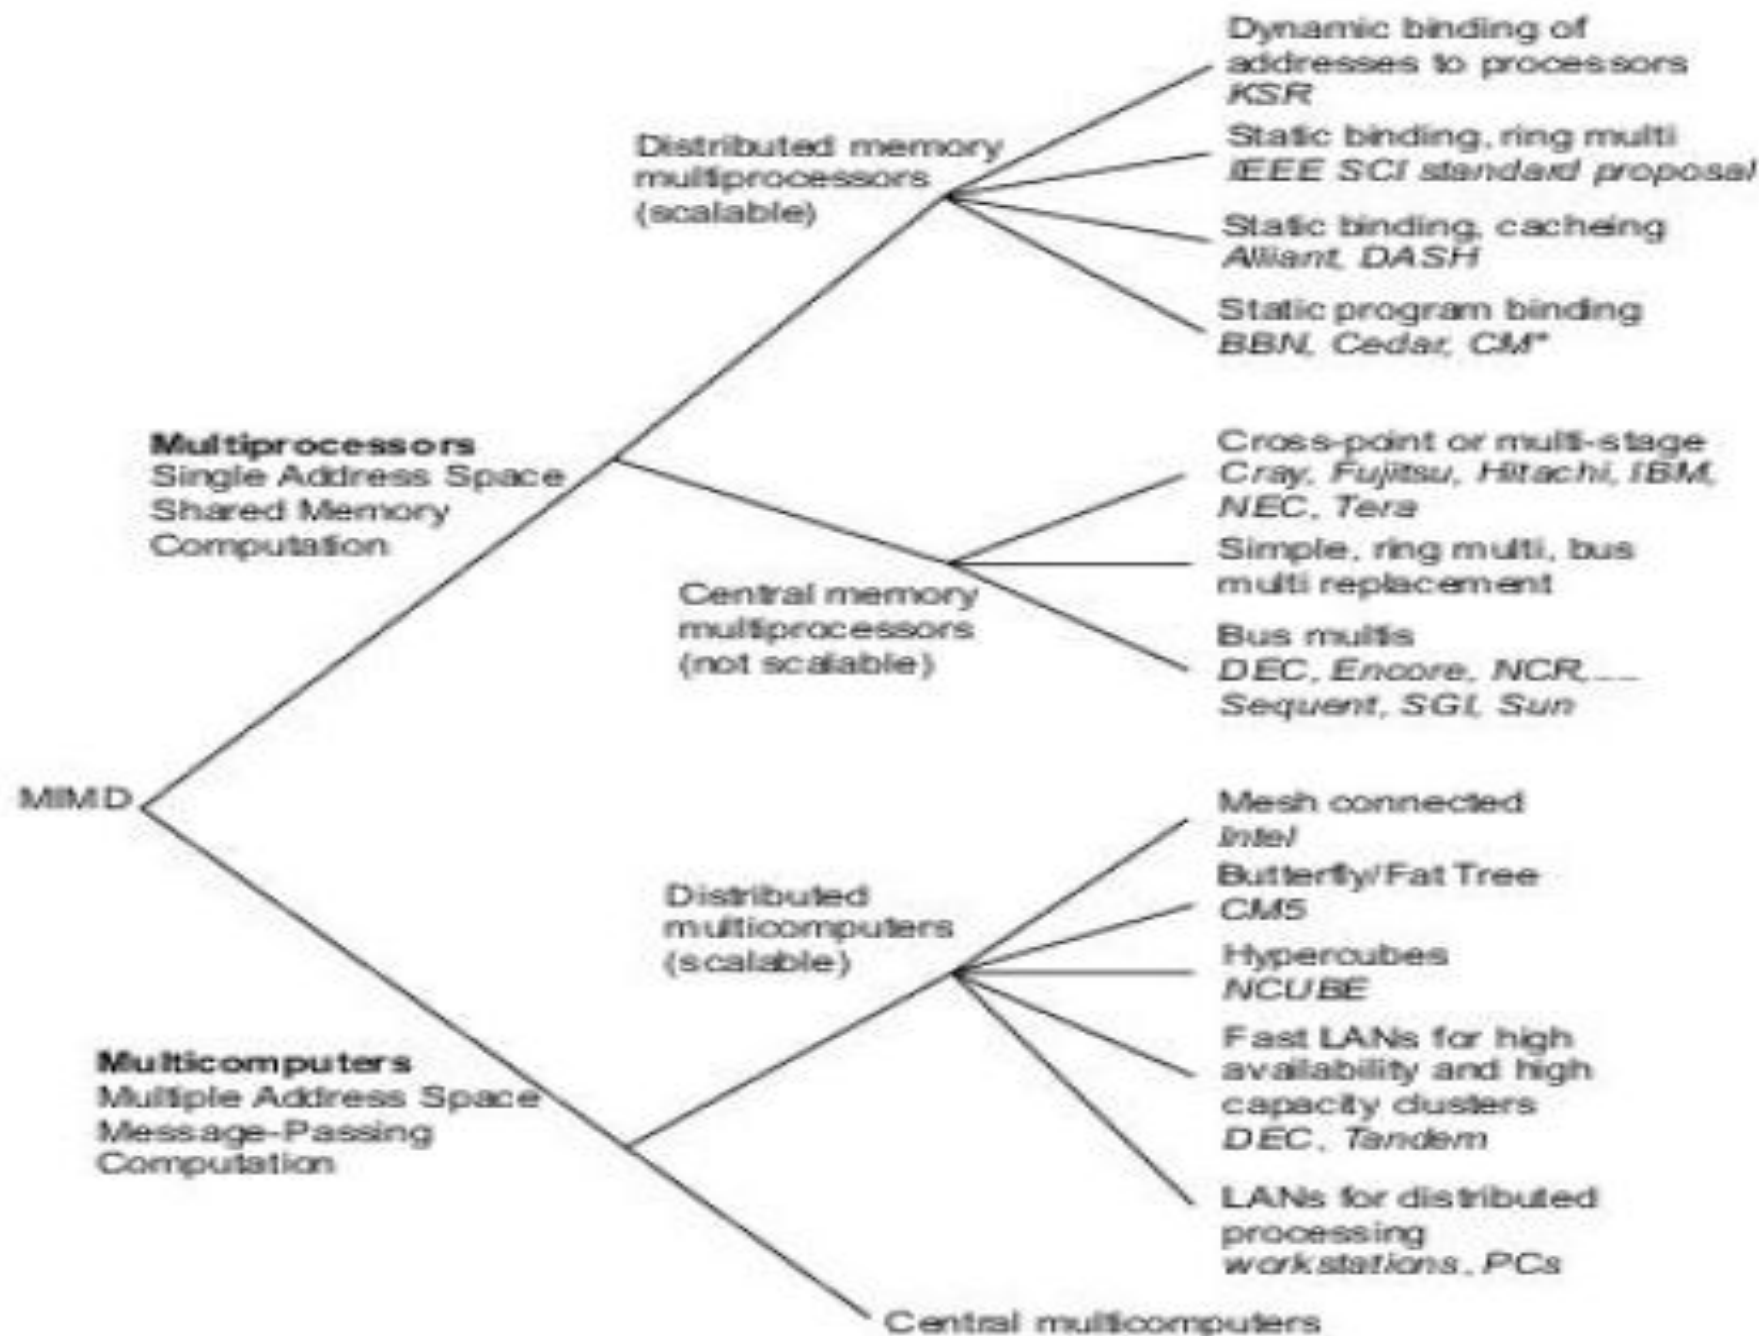
- System consists of multiple computers – nodes
- Interconnected by a message-passing network
- Node is an autonomous computer
- Consist of a processor, local memory, and disks or I/O peripherals
- Point-to-point static connections among the nodes
- Multicomputers also called no-remote-memory- access (NORMA)

- Accessible only by local processors



Generic model of a message passing multicomputers

<i>System Features</i>	<i>Intel Paragon XP/S</i>	<i>nCUBE/2 6480</i>	<i>Parsys Ltd. SuperNode1000</i>
Node Types and Memory	50 MHz i860 XP computing nodes with 16–128 Mbytes per node, special I/O service nodes.	Each node contains a CISC 64-bit CPU, with FPU, 14 DMA ports, with 1–64 Mbytes /node.	EC-funded Esprit supernode built with multiple T-800 Transputers per node.
Network and I/O	2-D mesh with SCSI, HIPPI, VME, Ethernet, and custom I/O.	13-dimensional hypercube of 8192 nodes, 512-Gbyte memory, 64 I/O boards.	Reconfigurable interconnect, expandable to have 1024 processors.
OS and Software Task Parallelism Support	OSF conformance with 4.3 BSD, visualization and programming support.	Vertex/OS or UNIX supporting message passing using wormhole routing.	IDRIS/OS UNIX-compatible.
Application Drivers	General sparse matrix methods, parallel data manipulation, strategic computing.	Scientific number crunching with scalar nodes, database processing.	Scientific and academic applications.
Performance Remarks	5–300 Gflops peak 64-bit results, 2.8–160 GIPS peak integer performance.	27 Gflops peak, 36 Gbytes/s I/O	200 MIPS to 13 GIPS peak.



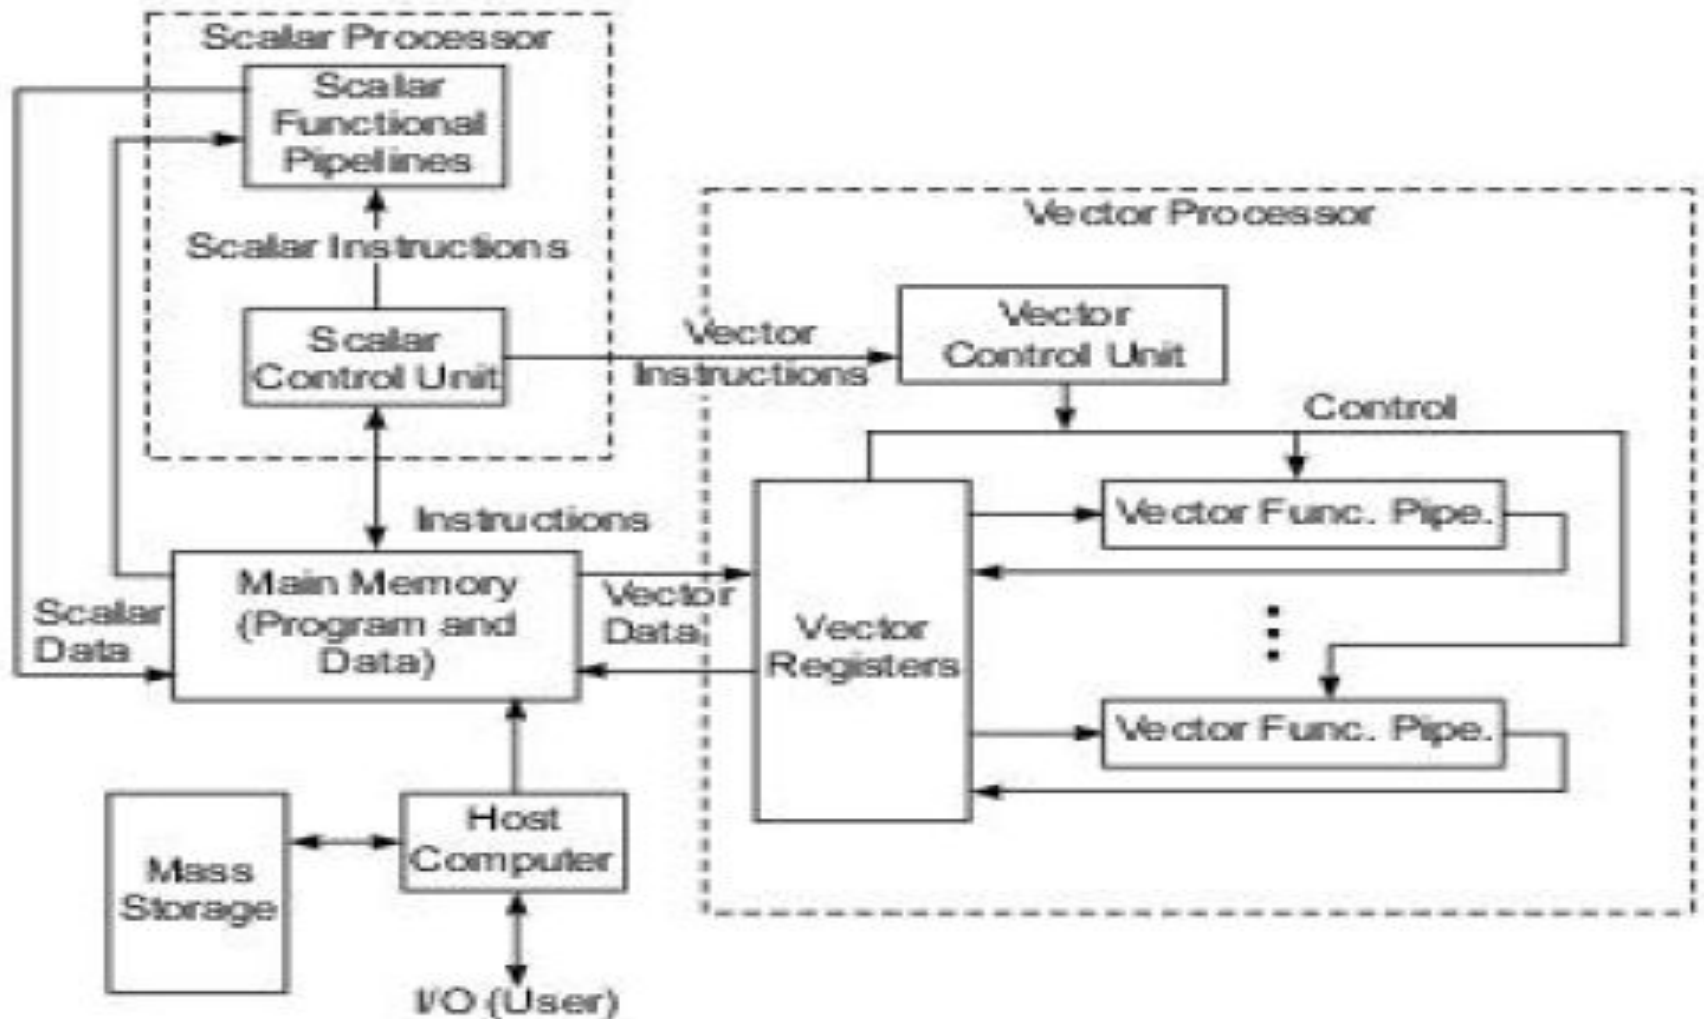
Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - **Multivector and SIMD computers**
 - Architectural development tracks
 - Conditions of parallelism

Multivector and SIMD computers

- Classify supercomputers either as
 - Vector Supercomputers
 - SIMD Supercomputers
- **Vector Supercomputer** – built on top of supercomputer
- Vector processor attached to scalar processor
- All instructions are first decoded by the scalar control unit
- If decoded instruction either a scalar operation or a program control operation : executed by the scalar processor using the scalar functional pipelines

- The architecture of a vector supercomputer

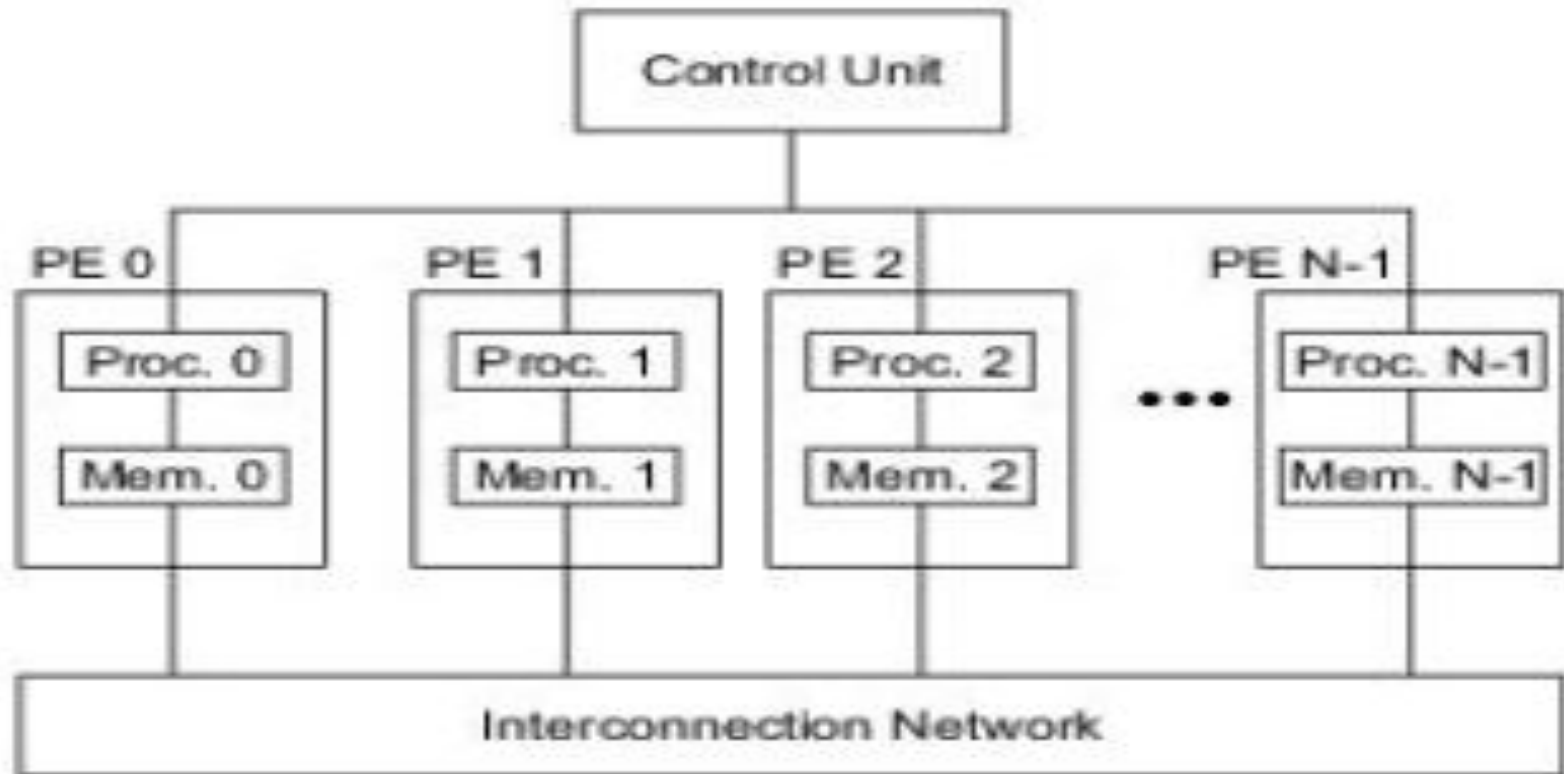


- Early Vector Processor Model:

<i>System Model</i>	<i>Vector Hardware Architecture and Capabilities</i>	<i>Compiler and Software Support</i>
Convex C3800 family	GaAs-based multiprocessor with 8 processors and 500-Mbyte/s access port. 4 Gbytes main memory. 2 Gflops peak performance with concurrent scalar/vector operations.	Advanced C, Fortran, and Ada vectorizing and parallelizing compilers. Also supported interprocedural optimization, POSIX 1003.1/OS plus I/O interfaces and visualization system
Digital VAX 9000 System	Integrated vector processing in the VAX environment, 125–500 Mflops peak performance. 63 vector instructions. $16 \times 64 \times 64$ vector registers. Pipeline chaining possible.	MS or ULTRIX/OS, VAX Fortran and VAX Vector Instruction Emulator (VVIEF) for vectorized program debugging.
Cray Research Y-MP and C-90	Y-MP ran with 2, 4, or 8 processors, 2.67 Gflop peak with Y-MP8256. C-90 had 2 vector pipes/CPU built with 10K gate ECL with 16 Gflops peak performance.	CF77 compiler for automatic vectorization, scalar optimization, and parallel processing. UNICOS improved from UNIX/V and Berkeley BSD/OS.

SIMD Supercomputers

- single instruction stream over multiple data streams



Operational model of SIMD computers

- An operational model of an SIMD computer

$$M=(N, C, I, M, R)$$

- {1} N is the number of processing elements (PEs) in the machine
- {2} C is the set of instructions directly executed by the control unit (CU)
- {3} I is the set of instructions broadcast by the CU to all PEs for parallel execution
- {4} M is the set of masking schemes
- {5} R is the set of data-routing functions

- Some Early Commercial SIMD Supercomputers

<i>System Model</i>	<i>SIMD Machine Architecture and Capabilities</i>	<i>Languages, Compilers and Software Support</i>
MasPar Computer Corporation MP-1 Family	Designed for configurations from 1024 to 16,384 processors with 26,000 MIPS or 1.3 Gflops. Each PE was a RISC processor, with 16 Kbytes local memory. An X-Net mesh plus a multistage crossbar interconnect.	Fortran 77, MasPar Fortran (MPF), and MasPar Parallel Application Language; UNIX/OS with X-window, symbolic debugger, visualizers and animators.
Thinking Machines Corporation, CM-2	A bit-slice array of up to 65,536 PEs arranged as a 10-dimensional hypercube with 4×4 mesh on each vertex, up to 1M bits of memory per PE, with optional FPU shared between blocks of 32 PEs, 28 Gflops peak and 5.6 Gflops sustained.	Driven by a host of VAX, Sun, or Symbolics 3600, Lisp compiler, Fortran 90, C*, and *Lisp supported by PARIS
Active Memory Technology DAP600 Family	A fine-grain, bit-slice SIMD array of up to 4096 PEs interconnected by a square mesh with 1 K bits per PE, orthogonal and 4-neighbor links, 20 GIPS and 560 Mflops peak performance.	Provided by host VAX/VMS or UNIX Fortran-plus or APAL on DAP, Fortran 77 or C on host.

Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - **Architectural development tracks**
 - Conditions of parallelism

Architectural Development Tracks

1. Multiple-Processor Tracks

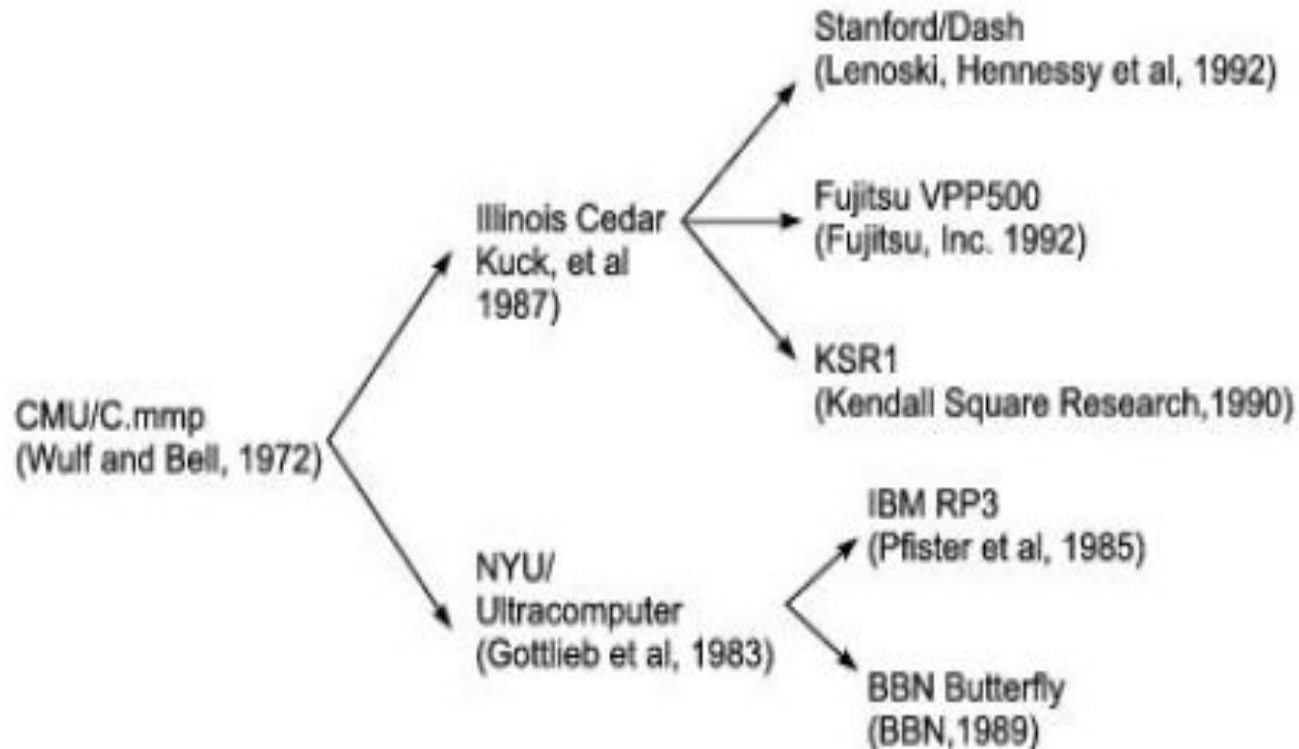
2. Multivector and SIMD Tracks

3. Multithreaded and Dataflow Tracks

1. Multiple-Processor Tracks(Shared-Memory Track and Message-passing Tracks)

- Multiple-processor system can be either a shared-memory multiprocessor or a distributed-memory multicomputer

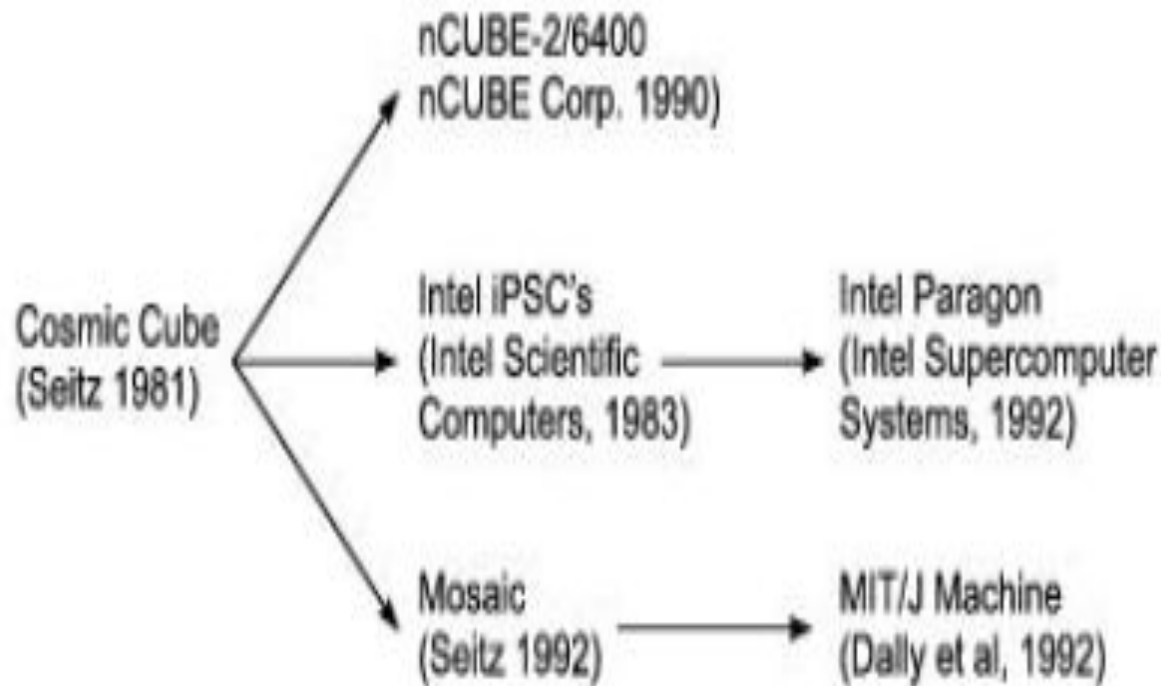
- **Shared-memory Track**



(a) Shared-memory track

- Employing a single address space in the entire system

- **Message-passing Track**

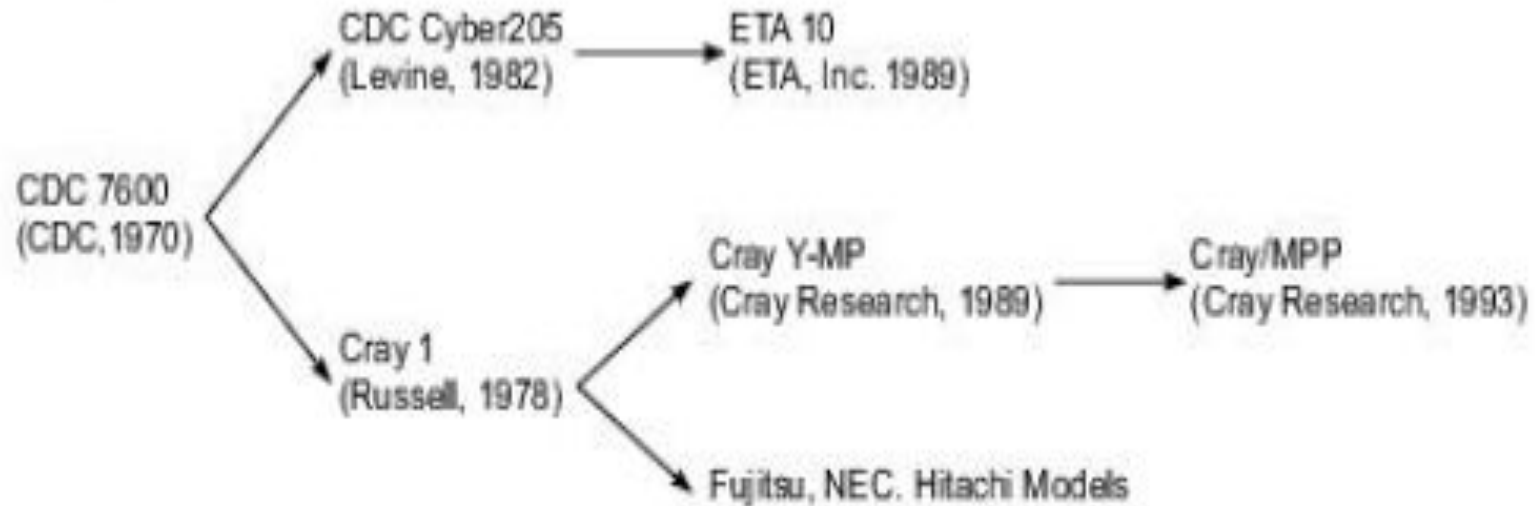


(b) Message-passing track

2. Multivector and SIMD Tracks

- **Multivector Track**

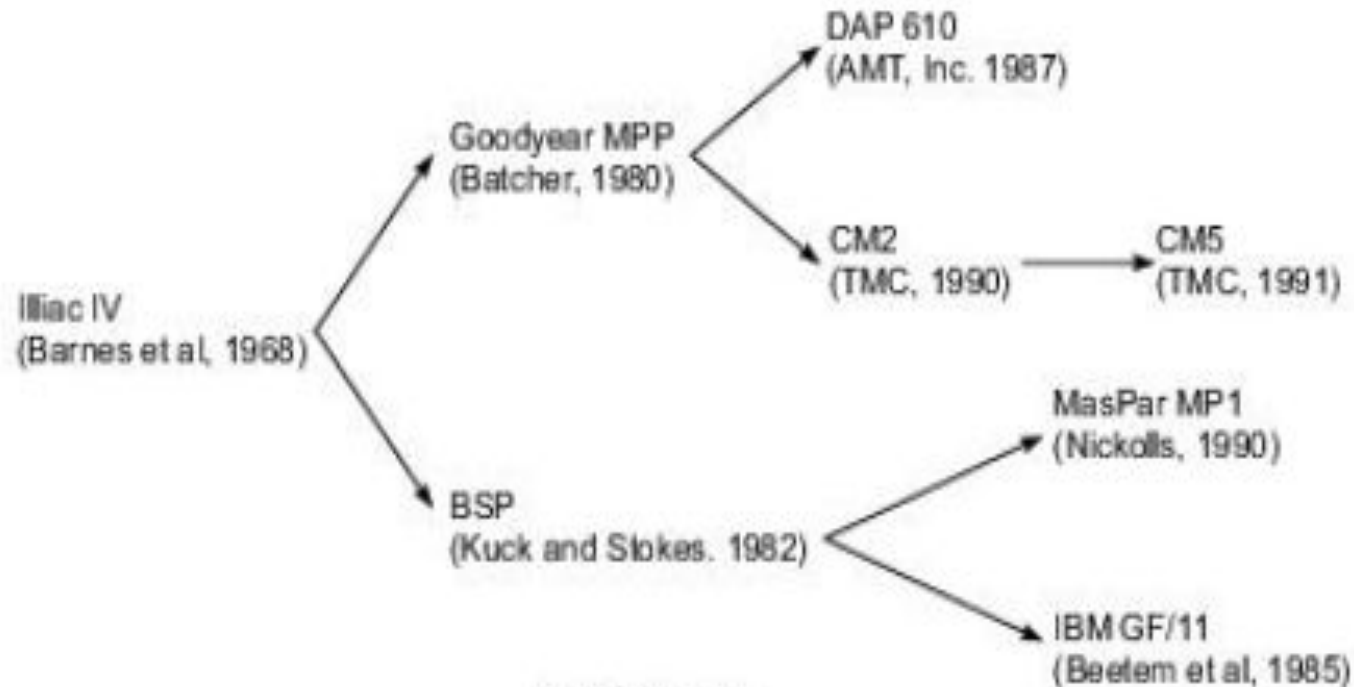
- Traditional vector supercomputers



(a) Multivector track


- **The SIMD Track**

- Illiac IV pioneered the construction of SIMD computers

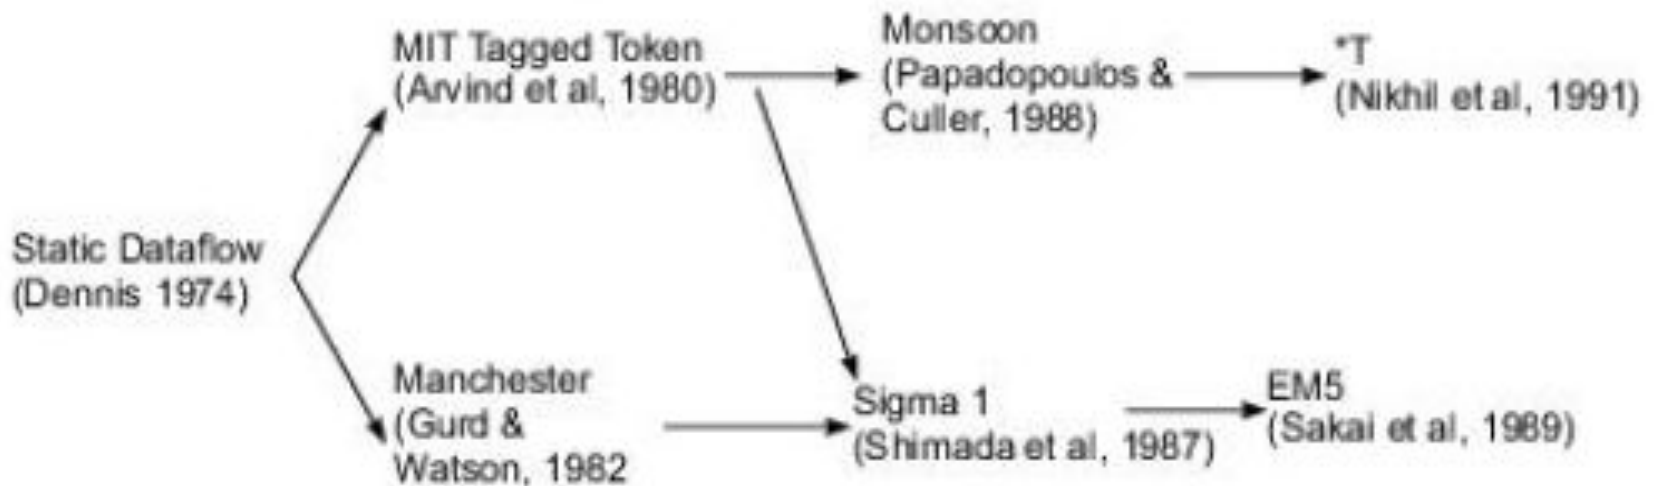


(b) SIMD track

3. Multithreaded and Dataflow Tracks



(a) Multithreaded track



(b) Dataflow track

- Each processor maintains a single thread of control with its hardware resources
- Multithreading - there are multiple threads of control in each processor
- Hiding long latency in building large-scale multiprocessors

- **The dataflow track**
 - To direct the program flow
 - Fine-grain instruction-level parallelism is exploited (**Fig(b)**)

Module 1

- **Parallel computer models**
 - Evolution of Computer Architecture
 - System Attributes to performance
 - Amdahl's law for a fixed workload
 - Multiprocessors and Multicomputers
 - Multivector and SIMD computers
 - Architectural development tracks
 - **Conditions of parallelism**

Conditions Of Parallelism

- To move parallel processing into the mainstream of computing make significant progress in
 - Computation Models
 - Inter Process Communications
 - System Integration

1) Data and Resource Dependence

- Dependence graph to describe the relations
 - **Data Dependence**

Five types of data dependence:

[1] *Flow dependence*: if at least one output of S1 feeds in as input to S2

$S1 \rightarrow S2.$

Eg: S1: $A = B + C$

S2: $D = F \times A$

1) Data and Resource Dependence

[2] *Antidependence*: if the output of S2 overlaps the input to S1

S1 \leftrightarrow S2

Eg: S1: $R1 = R3 + R7$

S2: $R3 = R4 + R5$

[3] *Output dependence*: Two statements are output dependence if they produce the same output variable.

$S1 \rightarrow S2$

Eg: S1: A = B + C

S2: A = E - F

[[4] *I/O dependence* : Read and write are I/O statements

S1 \Rightarrow S2

S1: Read (4),A(I)

S2: Process

S3: Write (4),B(I)


S4: Close(4)

[5] *Unknown dependence* : The dependence relation between two statements cannot be determined in the following situations:

- The subscript of a variable is itself subscribed.
- The subscript does not contain the loop index variable.
- A variable appears more than once with subscripts having different coefficients of the loop variable.
- The subscript is nonlinear in the loop index variable

Data dependence in programs

S1:	Load R1, A	/R1 \leftarrow Memory(A)/
S2:	Add R2, R1	/R2 \leftarrow (R1) + (R2)/
S3:	Move R1, R3	/R1 \leftarrow (R3)/
S4:	Store B, R1	/Memory(B) \leftarrow (R1)/



(a) Dependence graph

- **Control Dependence**

- Refers to the situation where the order of execution of statements cannot be determined before run time
- The successive iterations of the following loop are control-independent:

Control Dependence

```
      Do      101 = 1, N  
              IF (A(I - 1) .EQ. 0) A(I) = 0  
10  Continue
```

- **Resource Dependence**

- Concerned with the conflicts in using shared resources such as integer units, floating point units, registers, and memory areas, among parallel events
- When the conflicting resource is an ALU, call it ALU dependence

- **Bernstein's condition:**
 - Revealed a set of conditions based on which two processes can execute in parallel
- Define the input set i_i , of a process p_i , as the set of all input variables needed to execute the process

Bernstein's condition

- Consider two processes P1 and P2 with their input sets I1 and I2 and output sets O1 and O2 respectively
- 2 processes parallel $P1 \parallel P2$

Bernstein's condition

$$\left. \begin{array}{l} I_1 \cap O_2 = \phi \\ I_2 \cap O_1 = \phi \\ O_1 \cap O_2 = \phi \end{array} \right\}$$

- Three conditions are known as Bernstein's Conditions

DETECTION OF PARALLELISM

P1: $C = D \times E$

P2: $M = G + C$

P3: $A = B + C$

P4: $C = L + M$

P5: $F = G / E$

DETECTION OF PARALLELISM

P1: $C = D \times E$

P2: $M = G + C$

CONSIDER P1 AND P2

$$I_1 \cap O_2 = \emptyset$$

$$I_2 \cap O_1 = C \neq \emptyset$$

~~P1 || P2~~

DETECTION OF PARALLELISM

P1: $C = D \times E$

P3: $A = B + C$

CONSIDER P1 AND P3

$$I_1 \cap O_3 = \emptyset$$

$$I_3 \cap O_1 = C \neq \emptyset$$

~~P1 || P3~~

DETECTION OF PARALLELISM

$$P1: C = D \times E$$

$$P4: C = L + M$$

CONSIDER P1 AND P4

$$I_1 \cap O_4 = \emptyset$$

$$I_4 \cap O_1 = \emptyset$$

$$O_1 \cap O_4 = C \neq \emptyset$$

$$P1 \nparallel P4$$

DETECTION OF PARALLELISM

P1: $C = D \times E$

P5: $F = G/E$

CONSIDER P1 AND P5

$$I_1 \cap O_5 = \emptyset$$

$$I_5 \cap O_1 = \emptyset$$

$$O_1 \cap O_5 = \emptyset$$

$$P1 \parallel P5$$

DETECTION OF PARALLELISM

P2: $M = G + C$

P3: $A = B + C$

CONSIDER P2 AND P3

$$I_2 \cap O_3 = \emptyset$$

$$I_3 \cap O_2 = \emptyset$$

$$O_2 \cap O_3 = \emptyset$$

$$P2 \parallel P3$$

DETECTION OF PARALLELISM

P2: $M = G + C$

P4: $C = L + M$

CONSIDER P2 AND P4

$$I_2 \cap O_4 = C$$

~~P2~~ \parallel P3

DETECTION OF PARALLELISM

P2: $M = G + C$

P5: $F = G / E$

CONSIDER P2 AND P5

$$I_2 \cap O_5 = \emptyset$$

$$I_5 \cap O_2 = \emptyset$$

$$O_2 \cap O_5 = \emptyset$$

$$P2 \parallel P5$$

DETECTION OF PARALLELISM

P3: $A = B + C$

P4: $C = L + M$

CONSIDER P3 AND P4

$$I_3 \cap O_4 = C$$

~~P3 || P4~~

DETECTION OF PARALLELISM

P3: $A = B + C$

P5: $F = G / E$

CONSIDER P3 AND P5

$$I_3 \cap O_5 = \emptyset$$

$$I_5 \cap O_3 = \emptyset$$

$$O_3 \cap O_5 = \emptyset$$

$$P3 \parallel P5$$

DETECTION OF PARALLELISM

P4: $C = L + M$

P5: $F = G / E$

CONSIDER P3 AND P5

$$I_4 \cap O_5 = \emptyset$$

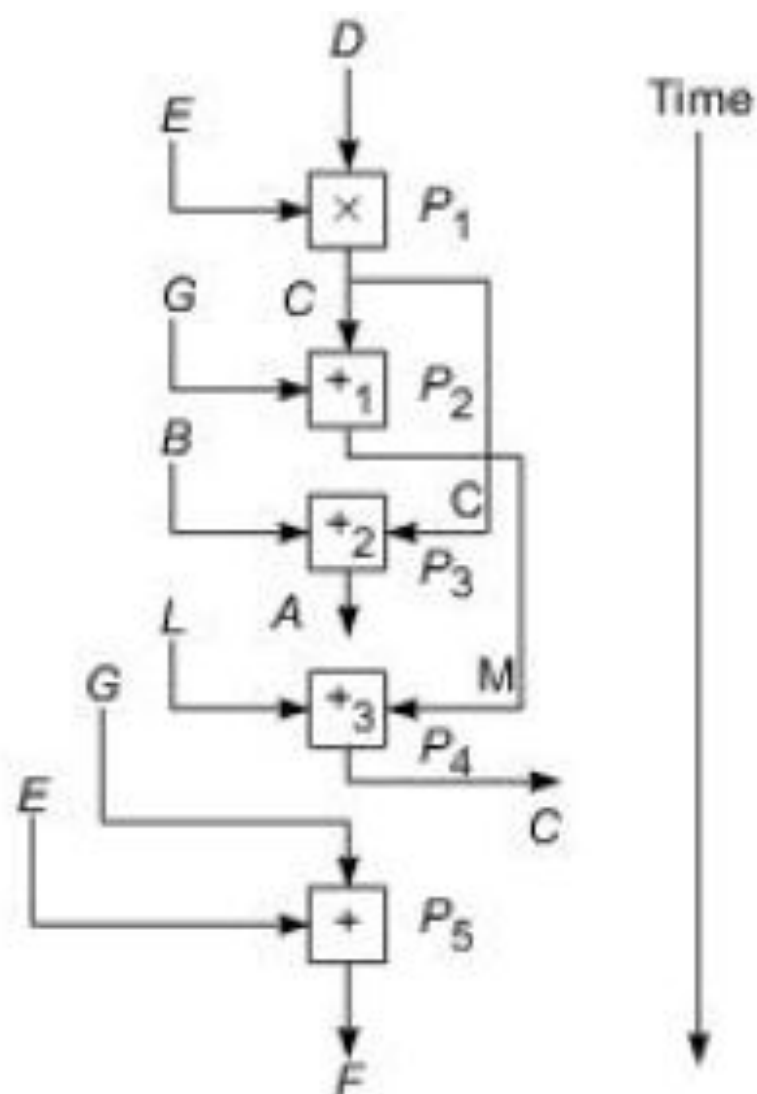
$$I_5 \cap O_4 = \emptyset$$

$$O_4 \cap O_5 = \emptyset$$


$$P4 \parallel P5$$

DETECTION OF PARALLELISM

- ONLY 5 pairs
- $P1 \parallel P5$, $P2 \parallel P3$, $P2 \parallel P5$, $P3 \parallel P5$ and $P4 \parallel P5$
can execute in parallel if there is no resource conflict
- Collectively $P2 \parallel P3 \parallel P5$



(b) Sequential execution in five steps, assuming one step per statement (no pipelining)



- (c) Parallel execution in three steps,
assuming two adders are available
per step

- **Hardware and Software Parallelism**

- Need special hardware and software support

- **Hardware Parallelism**

- Type of parallelism defined by the machine architecture and hardware multiplicity

- **Software Parallelism**

- Revealed in the program profile or in the program flow graph


- Software parallelism is it function of algorithm, programming style, and program design

Miss match

- $A = L1 * L2 + L3 * L4$
- $B = L1 * L2 - L3 * L4$

8 instruction –four load, 2 multi, 2 add/sub

- $8/3=2.67$




L_j : Load operation

X_j : Multiply operation

(a) Software parallelism

- 2-issue processor
- Execute one memory access(load & write) and one arithmetic (add,sub,mul)

- $8/7=1.14$



(b) Hardware parallelism