

# **MODULE III**

BY

Elizabeth Isaac

Department of Computer Science and Engg  
MA College of Engineering

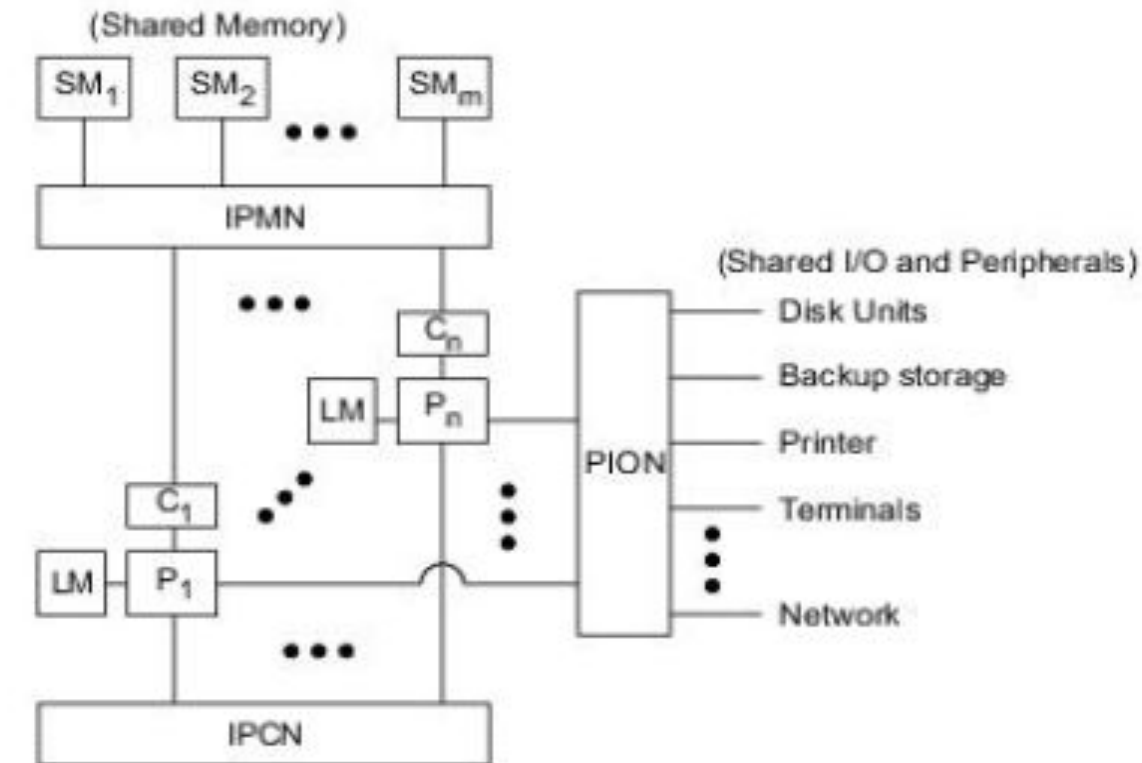
# MODULE III

- **Multiprocessors system interconnects**
  - Hierarchical bus systems
  - Cross bar switch and multiport memory
  - Multistage and combining networks
- Cache Coherence and Synchronization Mechanisms
  - Cache Coherence Problem
  - Snoopy Bus Protocol
  - Directory Based Protocol
  - Hardware Synchronization Problem

# **Multiprocessor System Interconnects**

- Parallel processing - efficient system interconnects for fast communication among multiple processors and shared memory, I/O, and peripheral devices
- Hierarchical buses, crossbar switches, and multistage networks

# Generalized Multiprocessor System



Legends:

- IPMN (Inter-Processor-Memory Network)
- PION (Processor- I/O Network)
- IPCN (Inter-Processor Communication Network)
- P (Processor)
- C (Cache)
- SM (Shared Memory)
- LM (Local Memory)

**Fig. 7.1** Interconnection structures in a generalized multiprocessor system with local memory, private caches, shared memory, and shared peripherals

- Network Characteristics

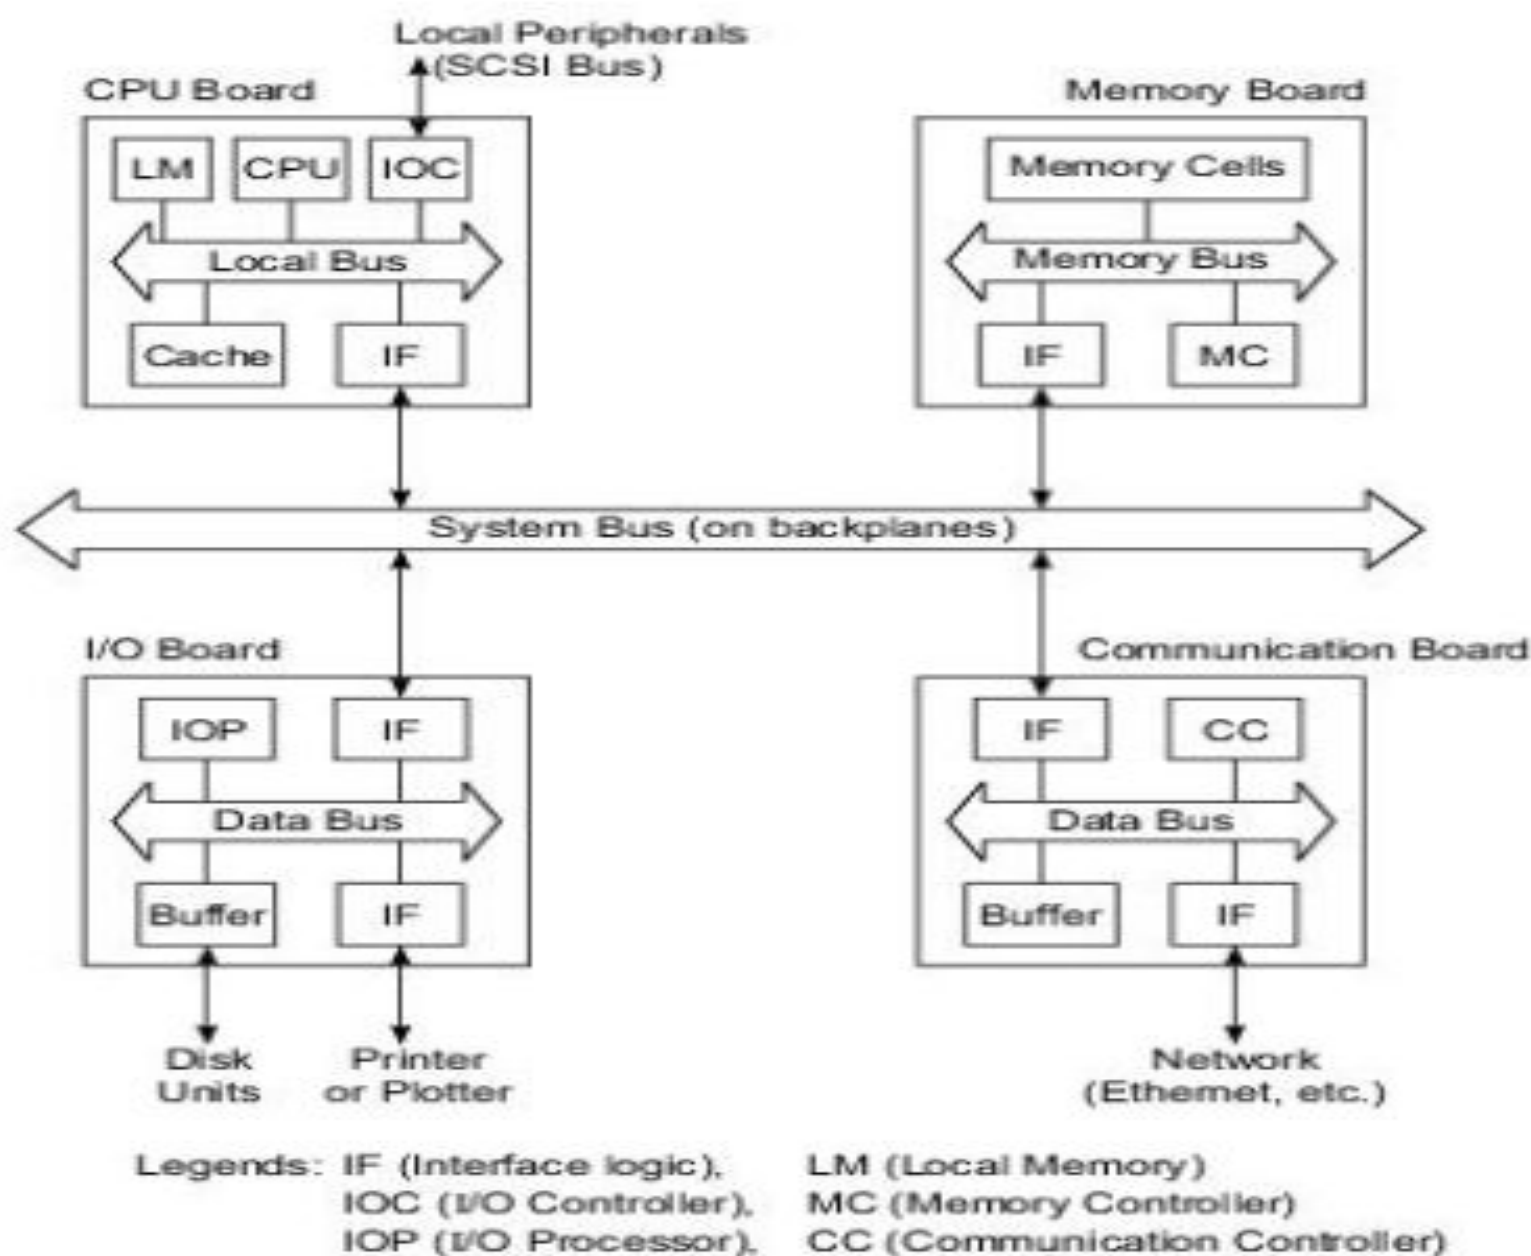
- Topology
  - **Dynamic Networks**
- Timing control protocol
  - **Synchronous** (with global clock)
  - **Asynchronous** (with handshake or interlocking mechanism)
- Switching method
  - **Circuit** switching
  - **Packet** switching
- Control Strategy
  - **Centralized** (global controller to receive requests from all devices and grant network access)
  - **Distributed** (requests handled by local devices independently)

# MODULE III

- Multiprocessors system interconnects
  - **Hierarchical bus systems**
    - Cross bar switch and multiport memory
    - Multistage and combining networks
- Cache Coherence and Synchronization Mechanisms
  - Cache Coherence Problem
  - Snoopy Bus Protocol
  - Directory Based Protocol
  - Hardware Synchronization Problem

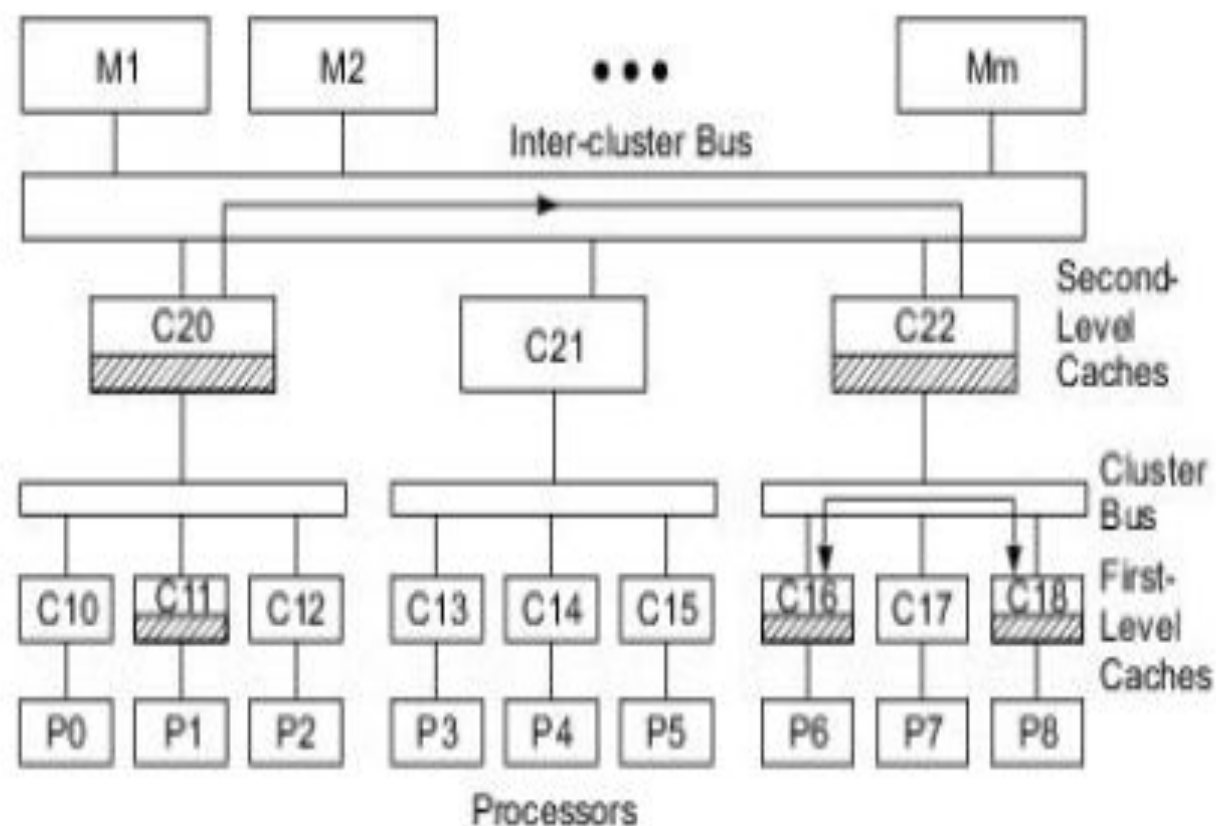
# **Hierarchical Bus Systems**

- Bus System: hierarchy of buses connecting various system and subsystem components in a computer
- Formed with a number of signal, control, and power lines
- Different buses - different interconnection functions



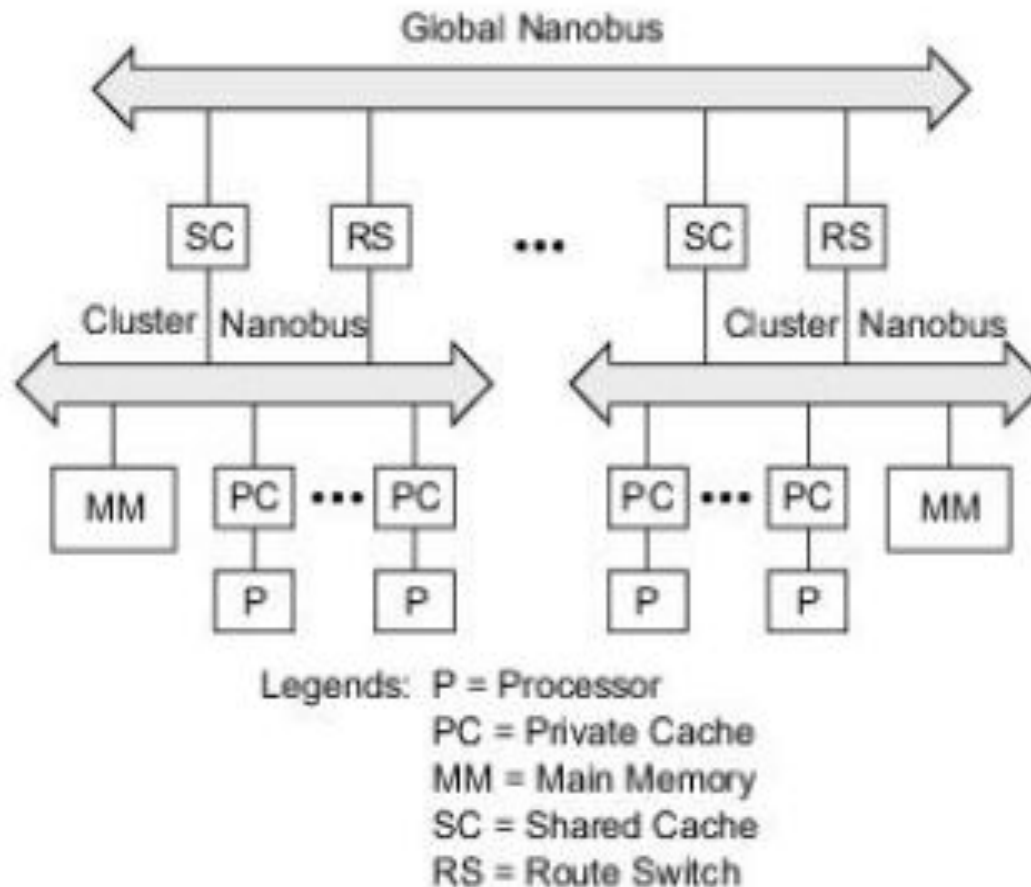
**Fig. 7.2** Bus systems at board level, backplane level, and I/O level





A hierarchical cache/bus architecture for designing a scalable multiprocessor (Courtesy of Wilson; reprinted from *Proc. of Annual Int. Symp. on Computer Architecture*, 1987)

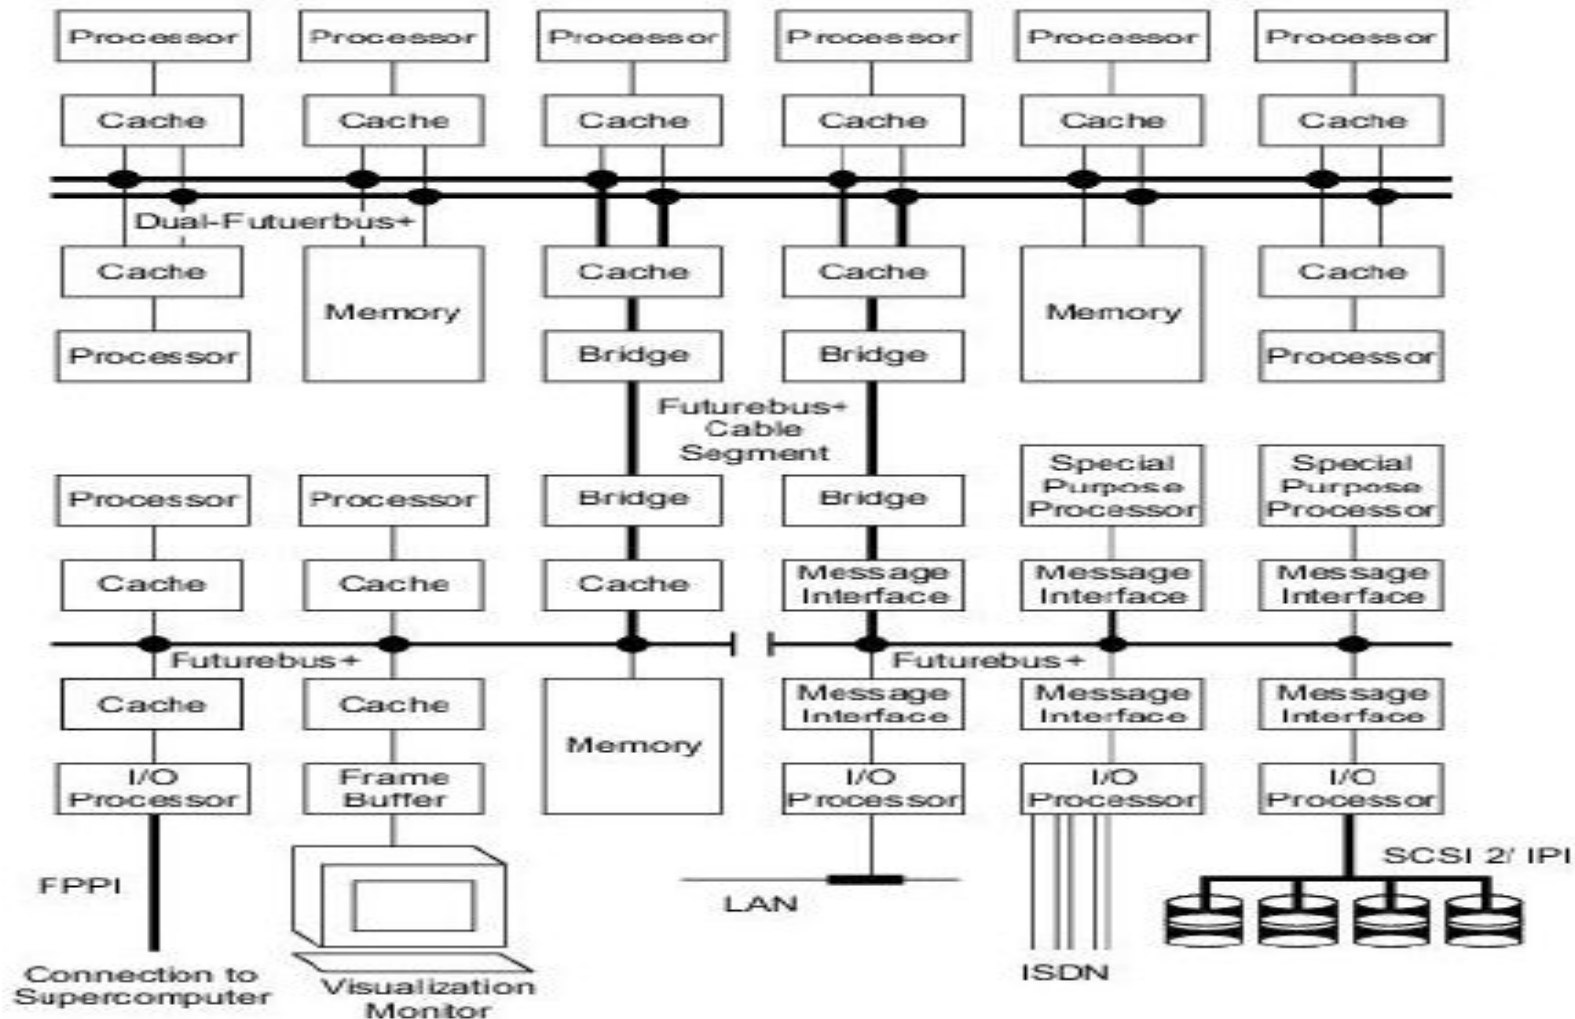
# Example 1.1 Encore Ultramax Multiprocessor Architecture



The Ultramax multiprocessor architecture using hierarchical buses with multiple clusters (Courtesy of Encore Computer Corporation, 1987)

- Two-level hierarchical-bus architecture
- The shared memories were distributed to all clusters instead of being connected to the intercluster bus
- When an access request reached the top bus, it would be routed down to the cluster memory that matched it with the reference address

- Idea of using bridges between multiprocessor clusters is to allow transactions initiated on a local bus to be completed on a remote bus



A multiprocessor system using multiple Futurebus+ segments (Reprinted with permission from IEEE Standard 896.1-1991, copyright © 1991 by IEEE, Inc.)

# MODULE III

- Multiprocessors system interconnects
  - Hierarchical bus systems
  - **Cross bar switch and multiport memory**
  - Multistage and combining networks
- Cache Coherence and Synchronization Mechanisms
  - Cache Coherence Problem
  - Snoopy Bus Protocol
  - Directory Based Protocol
  - Hardware Synchronization Problem

# Crossbar Switch and Multiport Memory

- Switched networks provide dynamic interconnections between the inputs and outputs
- **Network Stage:** Single-stage network is also called a recirculating network
  - cheaper to build, but multiple passes may be needed to establish certain connections


Multi-stage network

- **Blocking**
  - eg: Omega, Baseline and Banyan
- **Nonblocking Networks:**
  - Benes networks

- **Crossbar Networks** : every input port is connected to a free output port through a crosspoint switch without blocking
  - A crossbar network is a single-stage network built with unary switches at the crosspoints
  - Once the data is read from the memory. its value is returned to the requesting processor along the same crosspoint switch

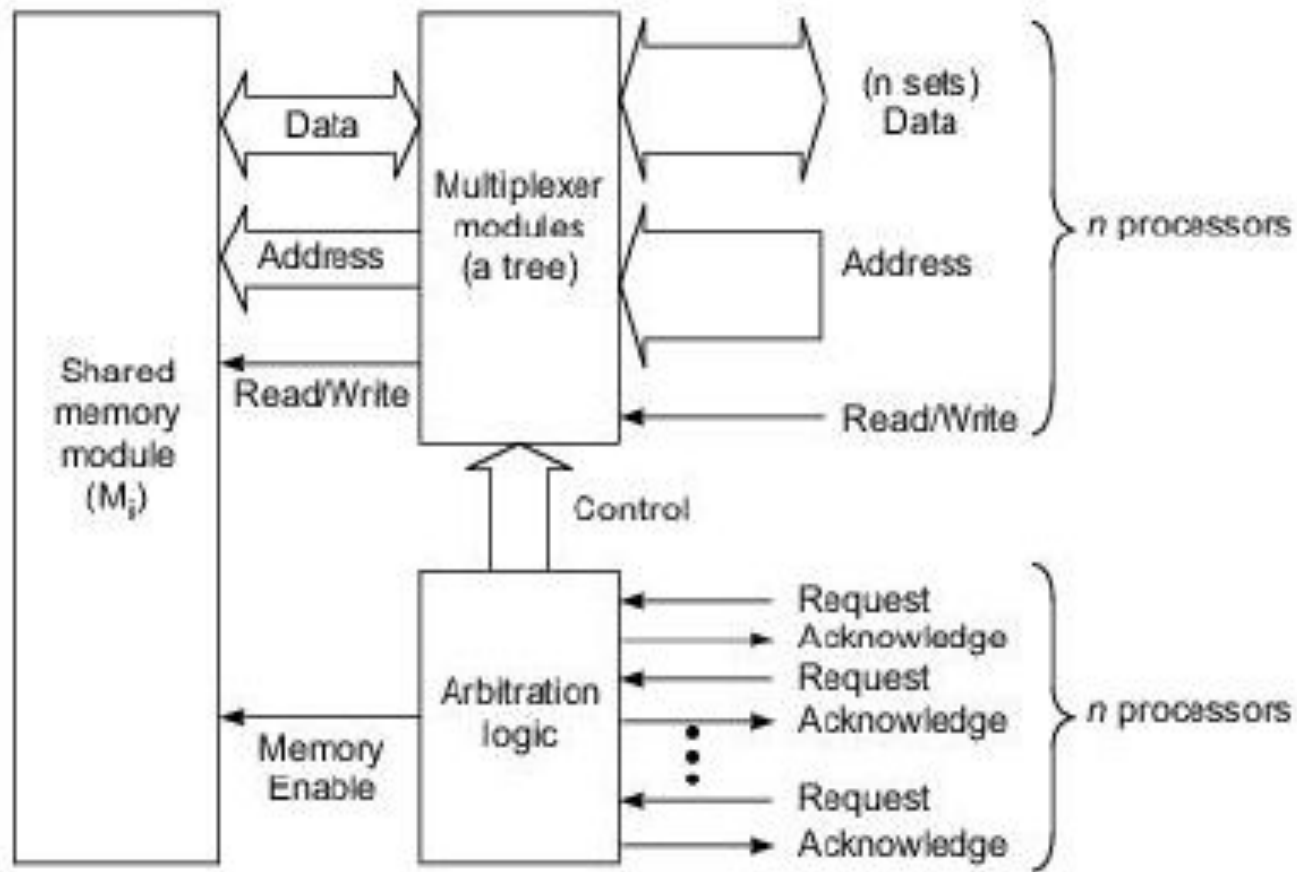


- **Crosspoint Switch Design** : Out of  $n$  crosspoint switches in each column of an  $n \times m$  or crossbar mesh, only one can be connected at a time
  - To resolve the contention for each memory module, each crosspoint switch must be designed with extra hardware

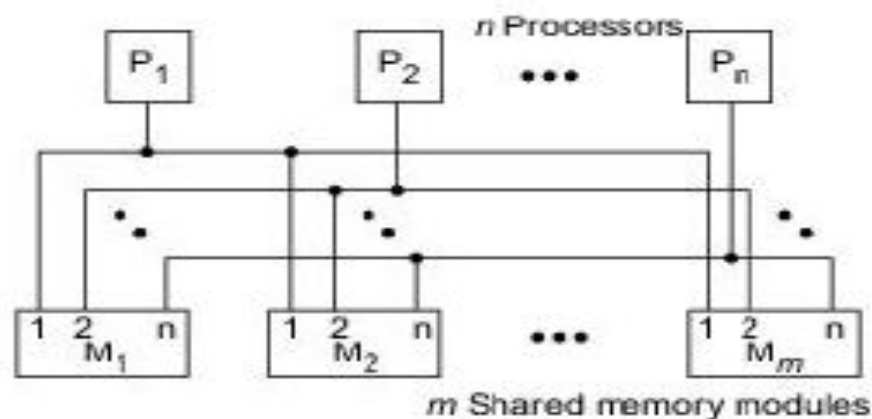


(a) Interprocessor-memory crossbar network built in the C.mmp multiprocessor at Carnegie-Mellon University (1972)

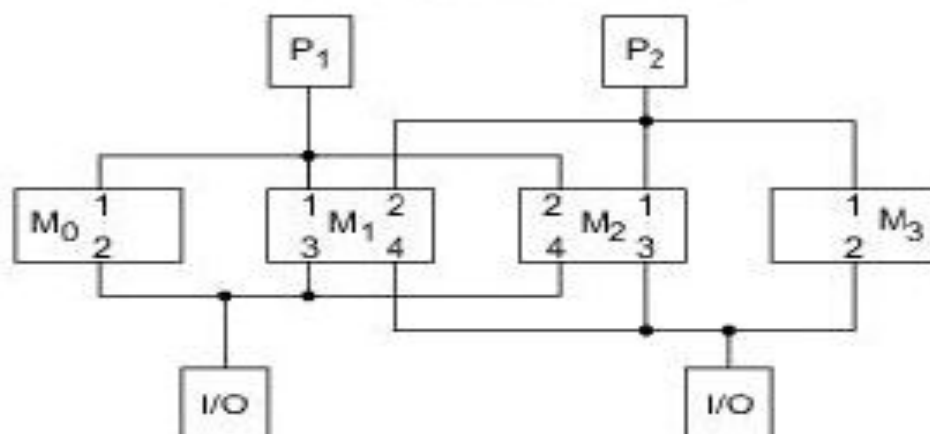
- Figure shows the schematic design of a row of crosspoint switches in a single crossbar network



- **Crossbar Limitations** : single processor can send many requests to multiple memory modules
  - cost-effective only for small multiprocessors with a few processors accessing a few memory modules
- **Multiport Memory** : memory module becomes more expensive due to the added access ports and associated logic.
  - Only one of  $n$  processor requests can be honored at a time



(a)  $n$ -port memory modules used



(b) Memory ports prioritized or privileged in each module by numbers


# MODULE III

- Multiprocessors system interconnects
  - Hierarchical bus systems
  - Cross bar switch and multiport memory
  - **Multistage and combining networks**
- Cache Coherence and Synchronization Mechanisms
  - Cache Coherence Problem
  - Snoopy Bus Protocol
  - Directory Based Protocol
  - Hardware Synchronization Problem

# Multistage and combining networks

- Multistage networks are used to build larger multiprocessor systems.
- Consists of more than one stage of switch boxes
- Perfect shuffle, butterfly, multiway shuffle , crossbar..
- multistage networks are the Omega network and the Butterfly network


# Multistage and combining networks




(e) 16x16 Omega network

# Properties - Multistage Interconnection Network (MIN).


- Consists of  $\log_2 N$  stages of  $2 \times 2$  switches
- Has  $N/2$  number of  $2 \times 2$  switches per stage
- Stages are labeled from 0 to  $\log_2 N - 1$
- There are 4 types of switch connection




(a) Straight



(b) Crossover



(c) Upper broadcast




(d) Lower broadcast




# Multistage and Combining Networks



**shuffle interconnection => rotate left**

$$S(a_{n-1} a_{n-2} \dots a_1 a_0) = (a_{n-2} a_{n-3} \dots a_0 a_{n-1})$$






(a) Perfect shuffle



000	----	>	000	----	>	000	----	>	000
001	----	>	010	----	>	100	----	>	001
010	----	>	100	----	>	001	----	>	010
011	----	>	110	----	>	101	----	>	011
100	----	>	001	----	>	010	----	>	100
101	----	>	011	----	>	110	----	>	101
110	----	>	101	----	>	011	----	>	110
111	----	>	111	----	>	111	----	>	111

## 8 x 8 OMEGA NETWORK




# Self Routing in Omega networks

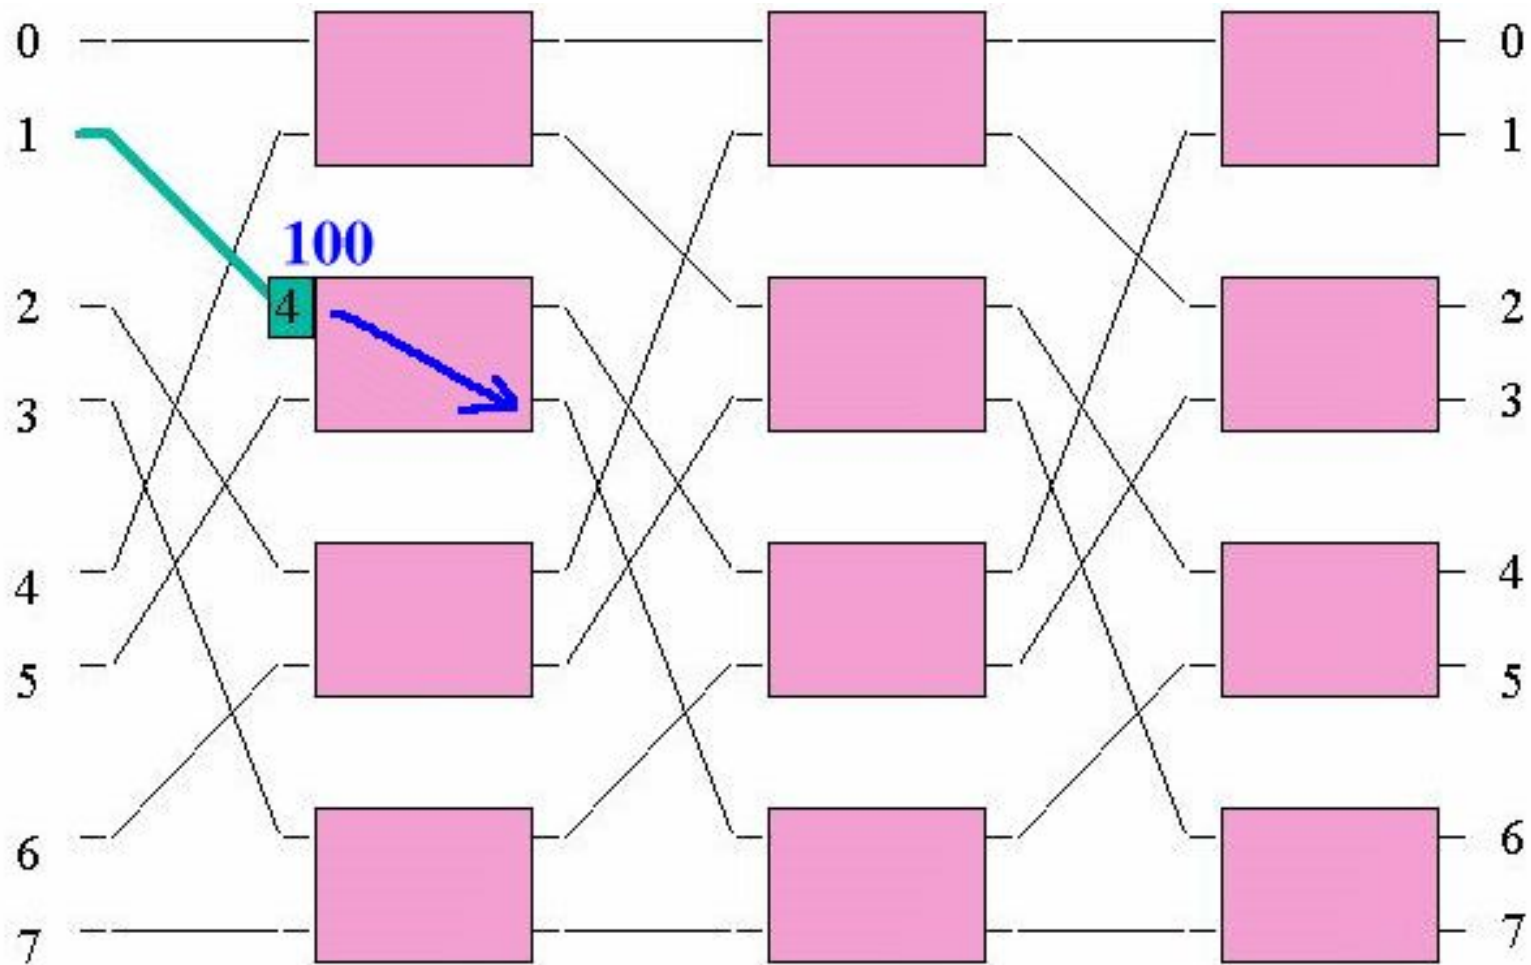
- Destination Tag Routing
- The source node generates a tag, which is binary equivalent of the destination.
- At each switch, the corresponding tag bit of destination is checked.
- If the bit is 0, the input is connected to the upper output.
- If bit is 1, the input is connected to the lower output.
- If both inputs have either 0 or 1, it is a switch conflict. One of them is connected. The other one is rejected or buffered at the switch


# Routing in Omega networks

- Suppose 1 to 4




# Routing in Omega networks





# Routing in Omega networks

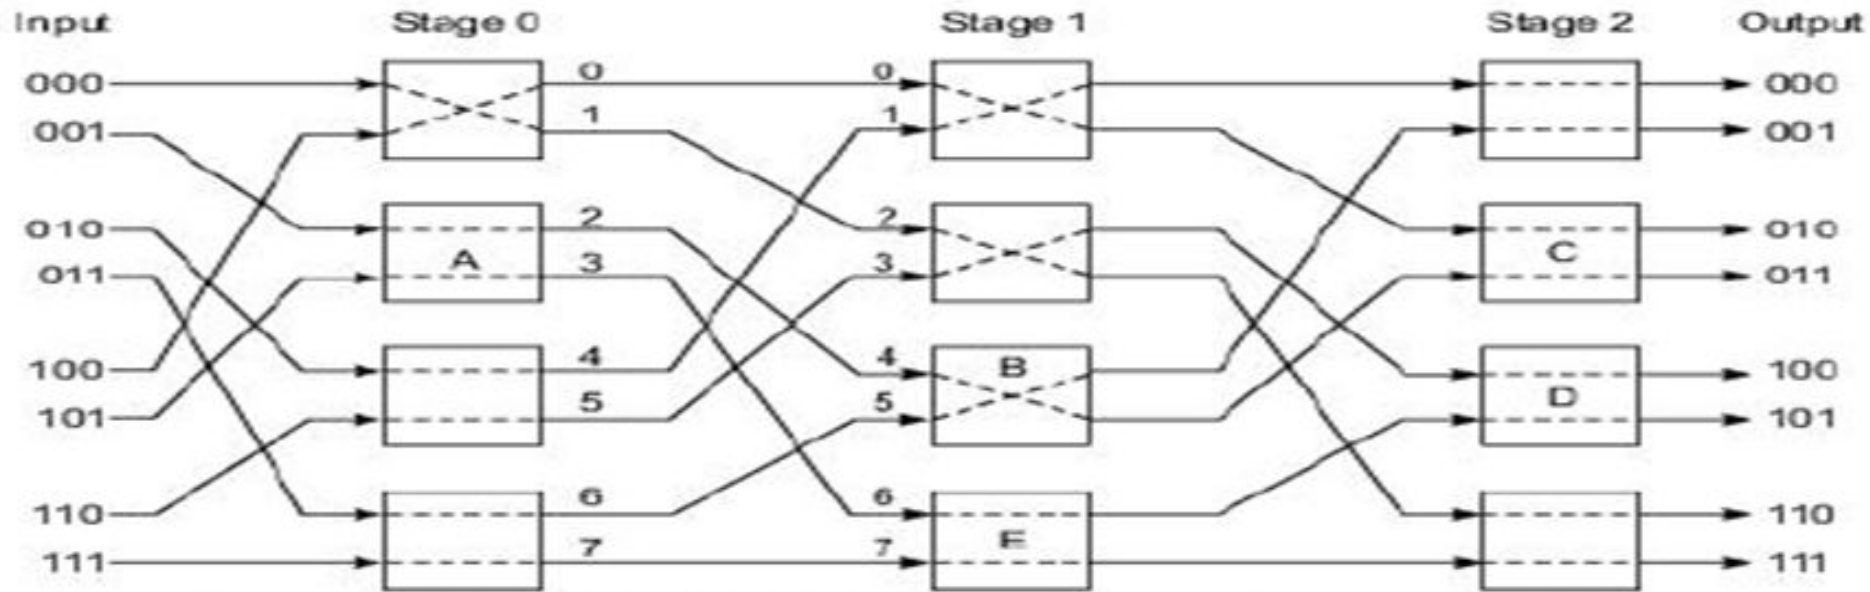




# Routing in Omega Networks


- permutations  $\Pi_1 = (0, 7, 6, 4, 2) (1, 3) (5)$
- $\Pi_2 = (0, 6, 4, 7, 3) (1, 5) (2)$  respectively

- Two switch setting of an 8\*8 Omega network built with 2\*2 switches




(a) Permutation  $\pi_1 = (0, 7, 6, 4, 2) (1, 3) (5)$  implemented on an Omega network without blocking


# Omega network



(b) Permutation  $\pi_2 = (0, 6, 4, 7, 3) (1, 5) (2)$  blocked at switches marked F, G, and H


- The Omega network can also-be used to broadcast data from one source to many destinations, as exemplified in Figure





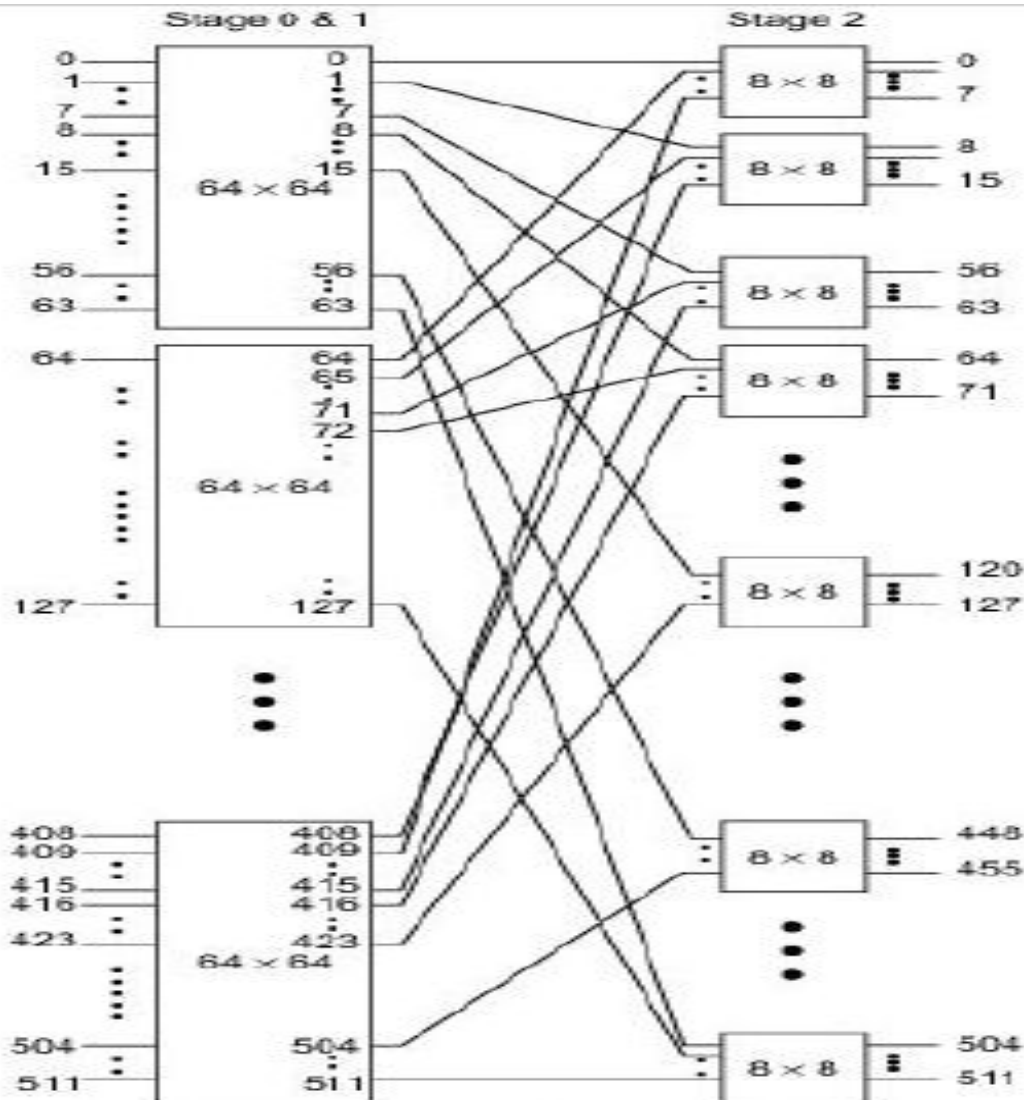
(b) Using four-way shuffle interstage connections

- **Routing In Butterfly Networks** : constructed with crossbar switches as building blocks



(a) A two-stage 64x64 Butterfly switch network built with 16 8x8 crossbar switches and eight-way shuffle interstage connections

# Modular construction of Butterfly switch networks with 8 X 8 crossbar switches



(b) A three-stage 512 x 512 Butterfly switch network built with 192 8 x 8 crossbar switches

# Fetch & add

- **The Hot-Spot Problem** : may appear corresponding to a certain memory module being excessively accessed by many processors at the same time the network performance significantly
- **Fetch & add** : This atomic memory operation is effective in implementing an N-way synchronization with a complexity independent of N.
  - In a Fetch & Add (x, e) operation , x is an
  - integer variable in shared memory
  - e is an integer increment



# Fetch & add

Fetch&Add ( $x, e$ )

$\{temp \leftarrow x;$   
 $x \leftarrow temp + e;$   
return  $temp\}$


# Fetch & add




(a) Two requests meet at a switch



(b) The switch forms the sum  $e_1 + e_2$ , stores  $e_1$  in buffer, and transmits the combined request to memory



(c) The original value stored in  $x$  is returned to switch, and the new value  $x + e_1 + e_2$  is stored in memory



(d) The values  $x$  and  $x + e_1$  are returned to  $P_1$  and  $P_2$  respectively

- Applications and Drawbacks
- Multistage Networks in Real System


# Cache Coherence And Synchronization Mechanisms

1. The Cache Coherence Problem
2. Directory-Based Protocols
3. Snoopy Bus Protocols
4. Hardware Synchronization Mechanisms

# The Cache Coherence Problem


- In a memory hierarchy for a multiprocessor system, data inconsistency may occur between adjacent levels or within the same level
- Caches in a multiprocessor environment introduce the Cache Coherence Problem
- Prevent the problem by maintaining a uniform state for each cached block of data

# Condition for Cache Incoherence



- This condition arises when the processor needs to share the writable data.
- In both policies **write back** and **write through**
- **Write back:** Write operations are usually made only to the cache. Main memory is only updated when the corresponding cache line is flushed from the cache.
- **Write through:** All write operations are made to main memory as well as to the cache, ensuring that main memory is always valid.

# Cache Coherence Problem Example



# CACHE COHERENCE & SYNCHRONIZATION MECHANISMS

- Cache protocols with the coherence for coping multi cache inconsistency problem
- Snoopy protocols are designed for bus-connected systems.
- Directory-based protocols apply to network-connected systems.
- Finally, hardware support for fast synchronization.



# Inconsistency in Data Sharing

In general, three sources of the problem are identified

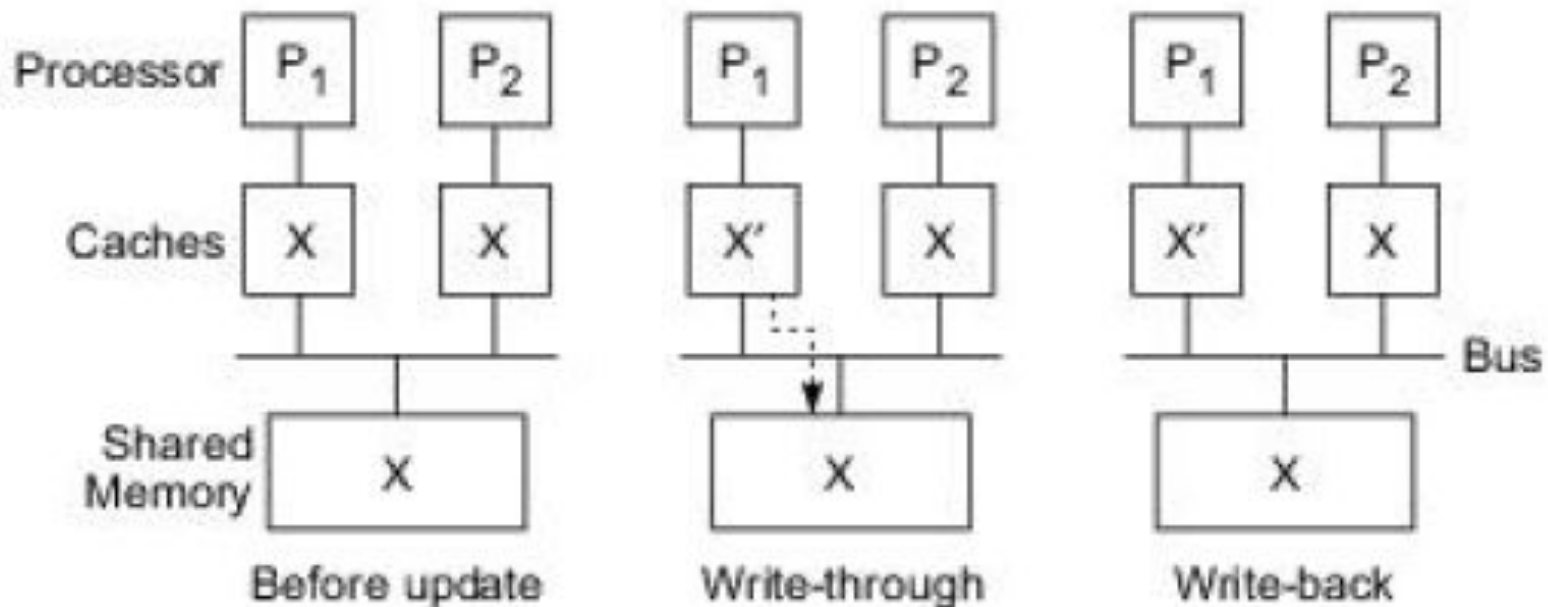
1. Inconsistency in Data sharing
2. Process Migration and I/O
3. Two Protocol Approaches

# 1. Inconsistency in Data sharing

- The cache inconsistency problem occurs only when **multiple private caches** are used.
- In general, three sources of the problem are identified:
  - sharing of writable data
  - process migration
  - I/O activity
- inconsistency may also occur when a **write-back** policy is used
- The main memory will be eventually updated when the modified data in the cache are replaced or invalidated

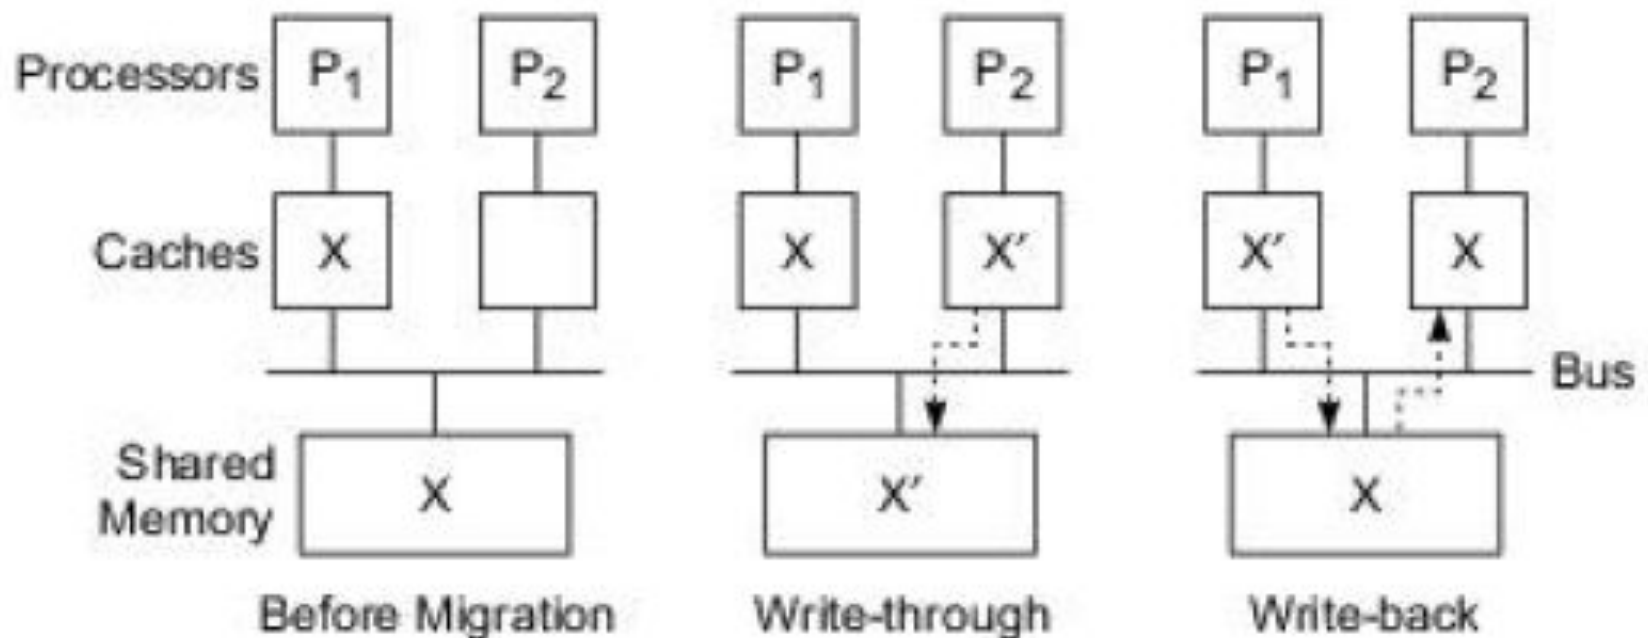
# 1. Inconsistency in Data sharing

- write-back policy
- Write-through policy




(a) Inconsistency in sharing of writable data

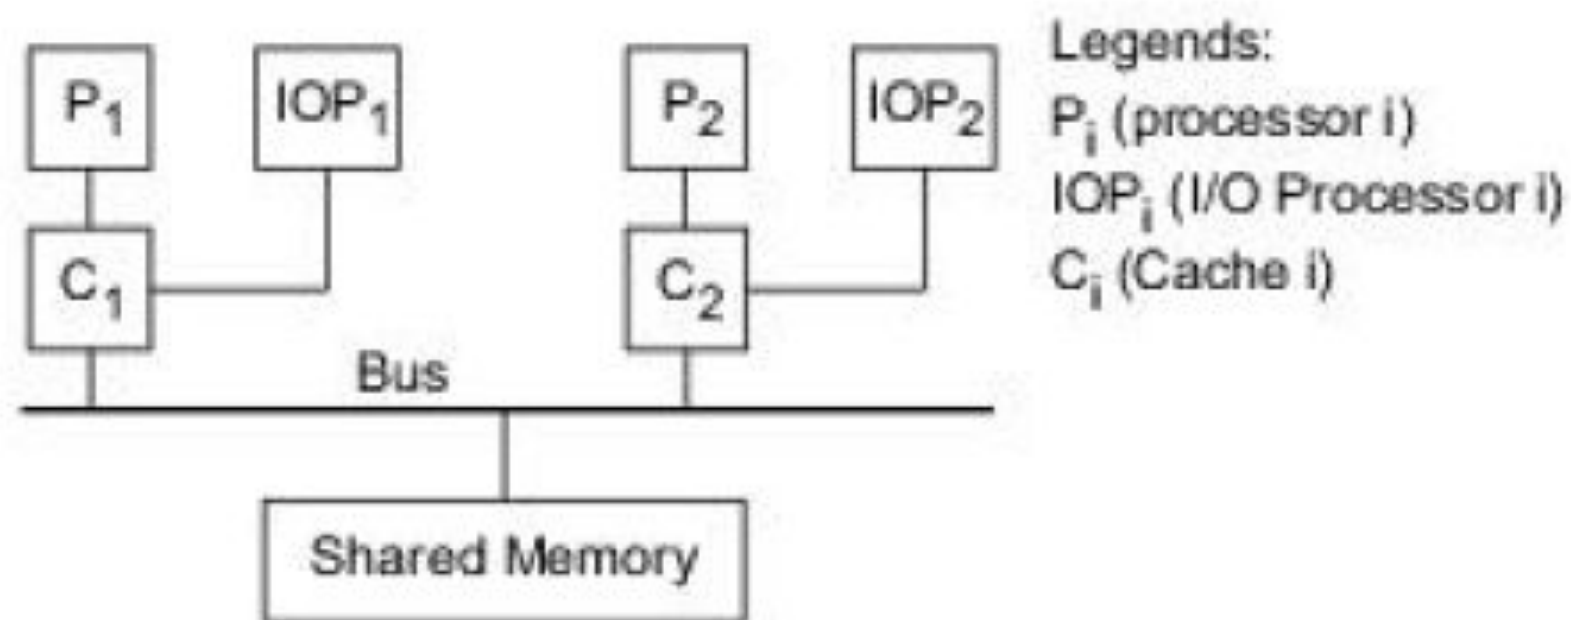
## 2. Process Migration and I/O



(b) Inconsistency after process migration



(a) I/O operations bypassing the cache



(b) A possible solution

- One possible **solution** to the I/O **inconsistency problem** is to
  - attach the **I/O processors** (IOP1, and IOP2) to the **private caches** (C1 and C1)

### 3. Two Protocol Approaches

- Many of the early commercially available multiprocessors used bus-based memory systems.
- A bus is a convenient device for ensuring cache coherence



# Directory-Based Protocols


1. Directory Structures
2. Full-Map Directories
3. Limited Directories
4. Chained Directories
5. Cache Design Alternative
6. Shared Cache
7. Non-cacheable Data
8. Cache Flushing

# 1. Directory Structures

- store information on where copies of cache blocks reside.
- Various directory-based protocols differ mainly in how the directory maintains information and what information it stores
- Central directory containing duplicates of all cache directories

# Directory Structures

- This central directory, providing
  - all the information needed to enforce consistency,
  - Is usually very large
  - must be associatively searched.
  - like the individual cache directories

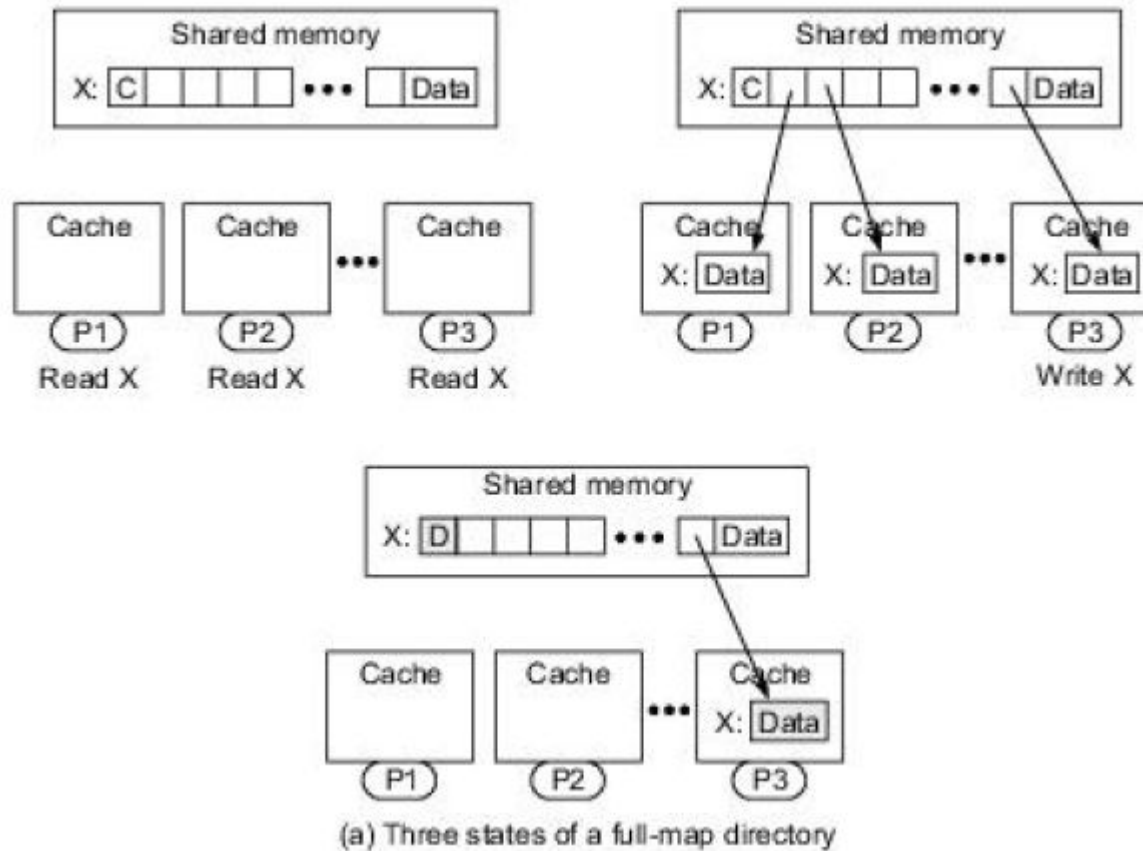


**Fig. 7.17** Basic concept of a directory-based cache coherence scheme (Courtesy of Censier and Feautrier, *IEEE Trans. Computers*, Dec. 1978)

## 2. Full-Map Directories

- implements directory entries with one bit per processor and a dirty bit
- Each bit represents the status of the block in the corresponding processor's cache
  - present or absent
- If the dirty bit is set,
  - then one and only one processor's bit is set and that processor can write into the block.

# Full-Map Directories



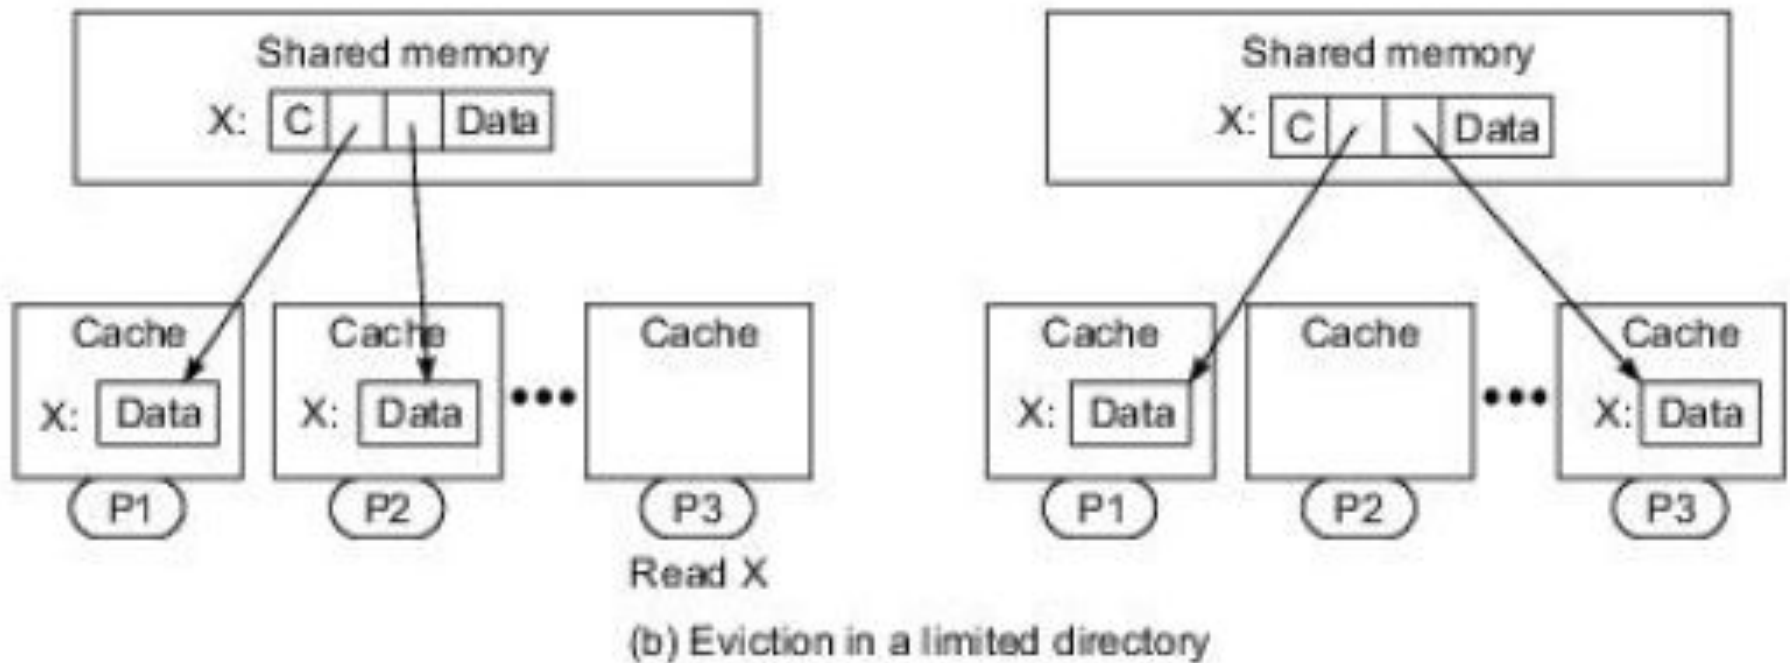
- In the **first state**, location X is missing in all of the caches in the system.
- The **second state** results from 3 caches requesting X.
- Three pointers (processor bits) are set

### 3. Limited Directories

- designed to solve the directory size problem
- Restricting the number of simultaneously cached copies of any particular block of data
- A directory protocol can be classified as  $Dir_i X$ 
  - $i$  is the number of pointers
  - $X$  is NB for a scheme with no broadcast
  - $Dir_N$  NB is full- map scheme without broadcast



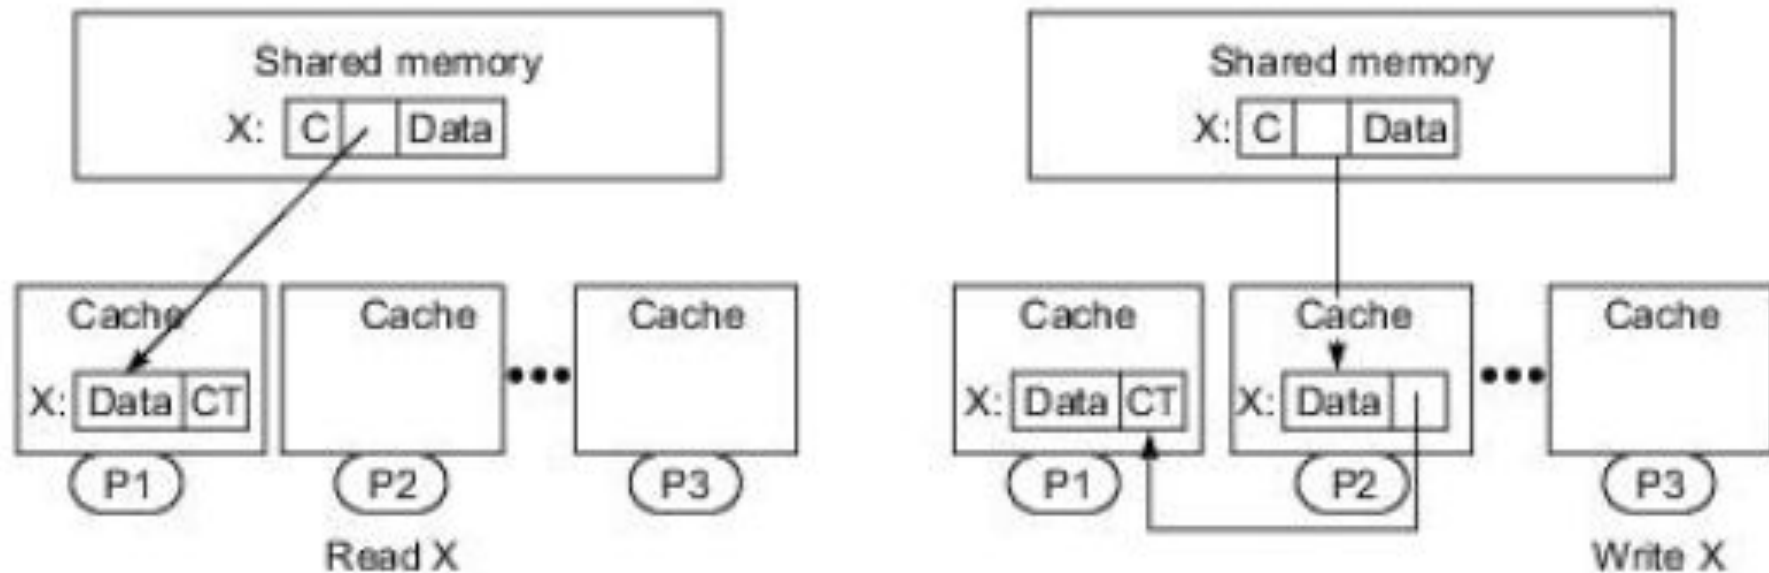
# Limited Directories



## 4. Chained Directories

- realize the scalability of limited directories without restricting the number of shared copies of data blocks.
- This type of cache coherence scheme is called a chained scheme
  - because it keeps track of shared copies of data by maintaining a chain of directory pointers

# Chained Directories



(c) The chained directory

# 5. Cache Design Alternative

- Each of the design alternatives has its own advantages and shortcomings
  - Shared Cache
  - Non-cacheable Data: The compiler must tag data as either Cacheable or non cacheable
  - Cache Flushing

# Hardware Synchronization Mechanisms

- Synchronization is a special form of communication in which control information is exchanged.
- use hardware mechanisms to implement low-level or primitive synchronization operations.
- or use software level synchronization mechanisms
- Semaphores or monitors

# Atomic Operation

- Most multiprocessors are equipped with hardware mechanisms for enforcing atomic operations
- **Operations-** memory read, write, or read-modify-write
- some interprocessor interrupts can be used for synchronization purposes.

- For example, the synchronization primitives, Test & Set (lock) and Reset (lock), are defined below:

Test & Set (lock)

temp  $\leftarrow$  lock; lock  $\leftarrow$  1;

return temp

Reset (lock)

lock  $\leftarrow$  0

- Test & Set is implemented with atomic read-modify-write memory operations.
- To synchronize concurrent processes, the software may repeat Test & Set until the returned value (temp) becomes 0.
- This synchronization primitive may tie up some bus cycles while a processor enters busy-waiting on the **spin lock**.

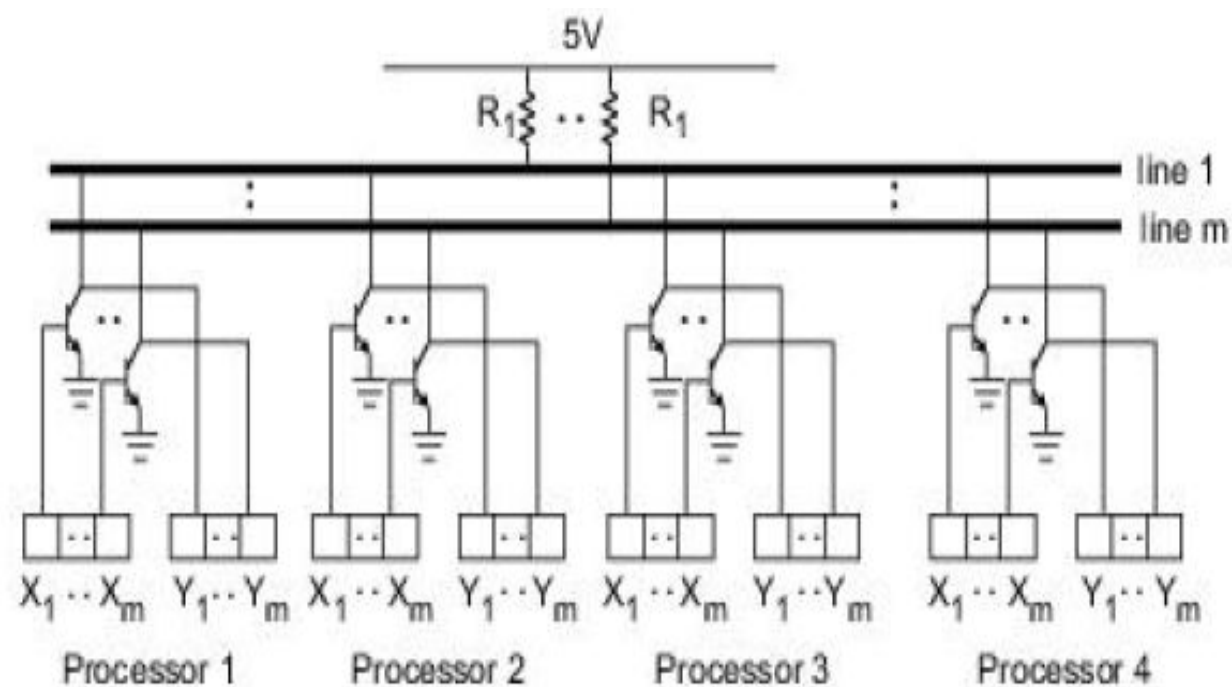


- When the lock is open, it signals all waiting processors through an interrupt.
- Concurrent processes residing in different processors can be synchronized using **barriers**

# Wired Barrier Synchronization

- A wired-NOR logic for implementing a barrier mechanism for fast synchronization
- control vector  $X = (X_1, X_2, \dots, X_m)$
- Monitor vector  $Y = (Y_1, Y_2, \dots, Y_m)$
- $m$ - barrier lines

- The number of barrier lines depends on the multiprogramming degree and the size of the multiprocessor system.
- Each **control bit  $X_i$** , is connected to the base (input) of a probing transistor
- The **monitor bit  $Y_i$**  checks the collector voltage (output) of the transistor.



(a) Barrier lines and interface logic

- Each barrier line is wired-NOR to  $n$  transistors from  $n$  processors.
- Whenever bit  $X_i$  is raised to high (1),
- the corresponding transistor is closed, pulling down (0) the level of barrier line  $i$ .
- The wired-NOR connection implies that line  $i$  will be high (1) only if control bits  $X_i$ , from all processors are low (0).

- The bit  $X_i$  is set to 1 when a process is initiated and reset to 0 when the process finishes its execution
- all processes finish their jobs,
  - the  $X_i$  bits from the participating processors are all set to 0
  - the barrier line is then raised to high (1)
  - signaling the synchronization barrier has been crossed

- only one barrier line is needed to monitor the initiation and completion of a single synchronization

Step 1: Forking (use of one barrier line)

	Processor 1	Processor 2	Processor 3	Processor 4
Line 1				
X	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
Y	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>

Step 2: Process 1 and Process 3 reach the synchronization point

X	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
Y	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
	Process 1	Process 2	Process 3	Process 4

Step 3: All processes reach the synchronization point


X	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
Y	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
	Process 1	Process 2	Process 3	Process 4

(b) Synchronization steps




# Snoopy Bus Protocols

- Write-invalidate
- Write-update




(a) Consistent copies of block X are in shared memory and three processor caches

# Snoopy Bus Protocols




(b) After a write-invalidate operation by  $P_1$

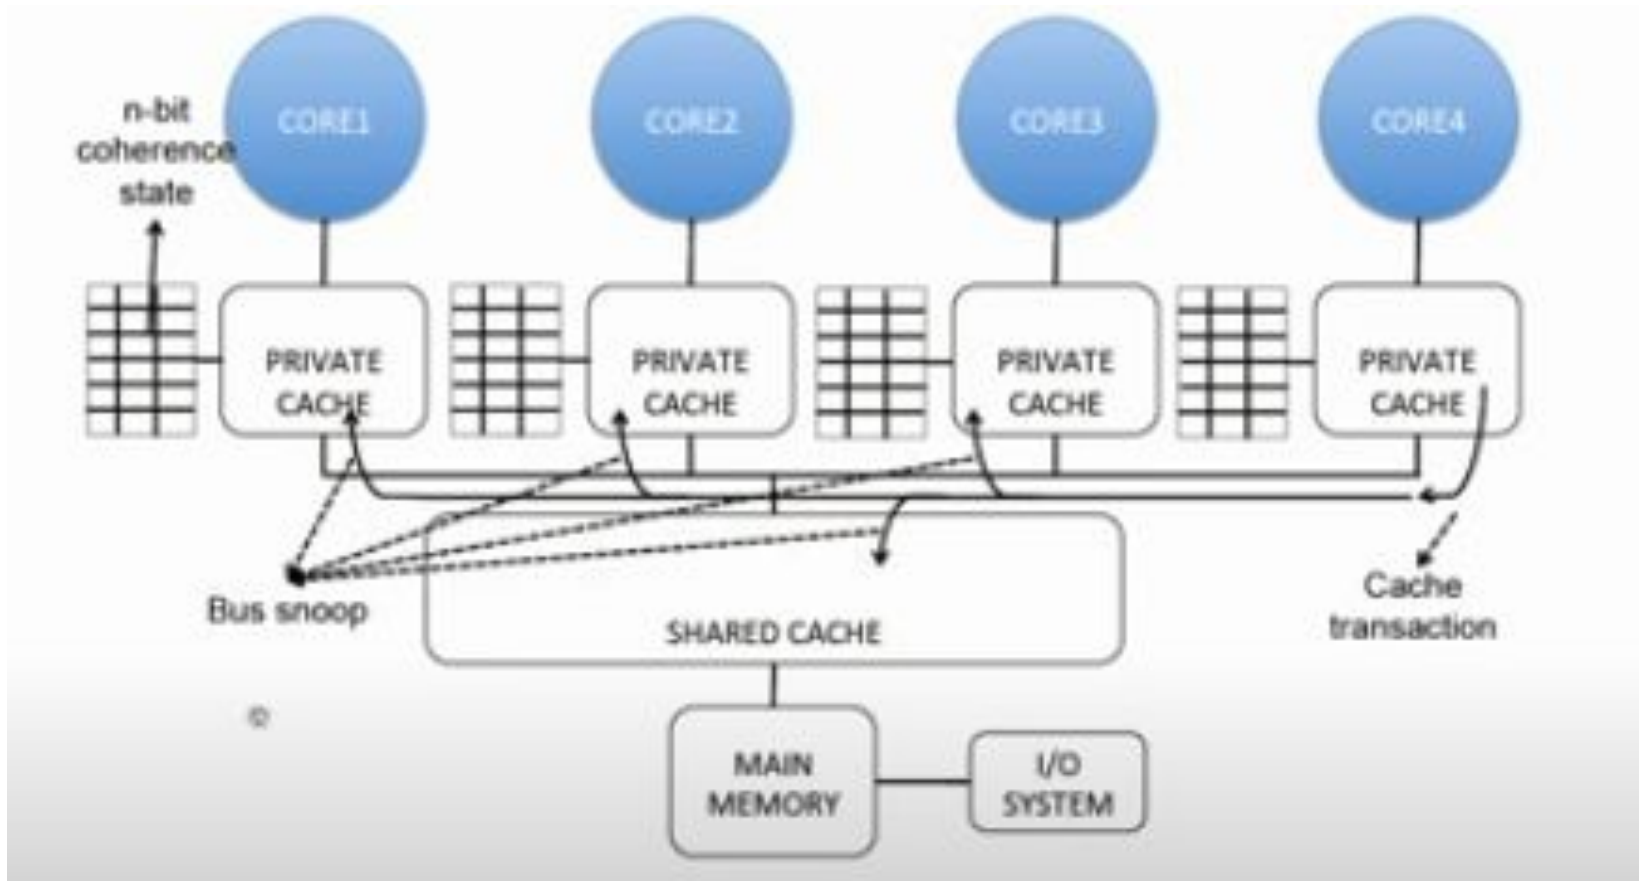


(c) After a write-update operation by  $P_1$

# Snoopy Bus Protocols



# Snoopy Bus Protocols




# Snoopy Bus Protocols

- Write-Back Caches
- Write-Through Caches
- Write-once Protocol

# Write-Through Caches

- The states of a cache block copy change with respect to **read** , **write** and **replacement** operations in the cache.
- A block copy of a write through cache i attached to processor can assume one of two possible cache states: **valid** or **invalid**



(a) Write-through cache

© 2009 ACPM, Inc. All rights reserved.

$W(i)$  = Write to block by processor  $i$ .

$R(i)$  = Read block by processor  $i$ .

$Z(i)$  = Replace block in cache  $i$ .

$W(j)$  = Write to block copy in cache  $j$  by processor  $j \neq i$ .

$R(j)$  = Read block copy in cache  $j$  by processor  $j \neq i$ .

$Z(j)$  = Replace block copy in cache  $j \neq i$ .

- In a *valid* state all processors can read ( $R(i)$ ,  $R(j)$ ) safely.
- Local processor  $i$  can also write ( $W(F1)$ ) safely in a *valid state*.
- The invalid state corresponds to the case of the block either being invalidated or being replaced ( $Z(i)$  or  $Z(j)$ ).



# Write-Back Caches


three-state coherence scheme

**RW** (Read-Write)

**RO** (Read-Only)

The INV (**invalidated or not-in-cache**) cache state

# Write-Back Caches



$W(i)$  = Write to block by processor  $i$ .  
 $R(i)$  = Read block by processor  $i$ .  
 $Z(i)$  = Replace block in cache  $i$ .

$W(j)$  = Write to block copy in cache  $j$  by processor  $j \neq i$ .  
 $R(j)$  = Read block copy in cache  $j$  by processor  $j \neq i$ .  
 $Z(j)$  = Replace block copy in cache  $j \neq i$ .

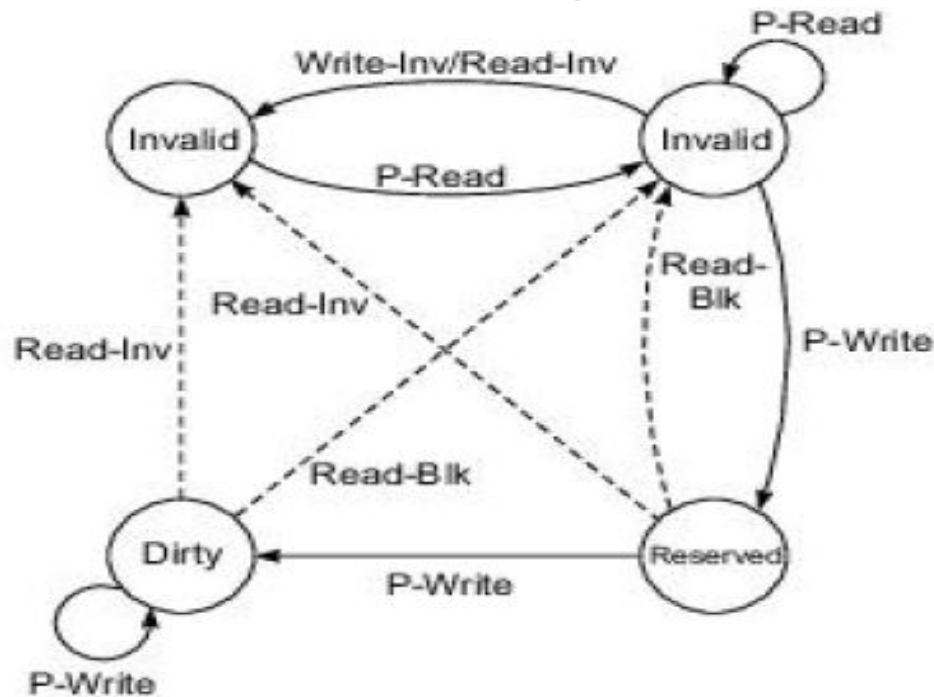
(b) Write-back cache

State-transition graph for a cache block using write-invalidate snoop protocol

# Write-once Protocol

- combines the advantages of both **write-through** and **write-back invalidations**.
- In order to reduce bus traffic, the very first **write** of a cache block uses a write-through policy.
- This will **result** in a consistent memory copy while all other cache copies are invalidated.
- After the first write, shared memory is updated using a **write—back policy**.

- Four cache states are given here:



Solid lines: Command issued by local processor  
 Dashed lines: Commands issued by remote processors  
 via the system bus.

Goodman's write-once cache coherence protocol using the write invalidate policy on write-back caches (Adapted from James Goodman 1983, reprinted from Stenstrom, *IEEE Computer*, June 1990)

- **Valid** :The cache block, which is consistent with the memory copy, has been read from shared memory and has not been modified.
- **Invalid** :The block is not found in the cache or is inconsistent with the memory copy.
- **Reserved** :Data has been written exactly once since being read' from shared memory.
  - The cache copy is consistent with the memory copy, which is the only other copy.

- **Dirty**: The cache block has been modified (written) more than once, and the cache copy is the only one in the system
  - thus inconsistent with all other copies

## 4.Cache Event and Action

- The memory-access and invalidation commands trigger the following events and actions:
  - **Read-miss**: When a processor wants to read a block that is not in the cache
  - **Write-hit**: If the copy is in the dirty or reserved state, the write can be carried out locally

- **Write-miss:** When a processor Fails to write in a local cache, the copy must come either from the main memory or from a remote cache with a dirty block.
- **Read-hit:** Read-hits can always be performed in a local cache without causing a state transition or using the snoopy bus for invalidation.
- **Block Replacement.**‘ If a copy is dirty it has to be written back to main memory by block replacement.



# Multilevel Cache Coherence

- To maintain consistency among cache copies at various levels
- Consistency of caches at different levels

# Protocol Performance issue

- depends heavily on the workload patterns and implementation efficiency.
  - is to reduce bus traffic,
  - reducing the effective memory-access time
- The block size is **very sensitive**

- For a uniprocessor system
  - bus traffic and memory-access time are mainly contributed by cache misses.
- For larger caches
  - the data pollution point appears at a larger block size.
- For a system requiring extensive process migration or synchronization.
- the write-invalidate protocol will perform better