

DISTRIBUTED COMPUTING

Module VI

COORDINATION AND AGREEMENT

DISTRIBUTED MUTUAL EXCLUSION

- Distributed processes often need to coordinate their activities
- Mutual exclusion is required to prevent interference and ensure consistency when accessing the shared resources
- This is the ***critical section problem*** (*familiar in operating systems*)
- A solution to *distributed mutual exclusion is required: based solely on message passing*

Algorithms for mutual exclusion

- Consider a system of N processes $p_i, i = 1, 2, \dots, N$, that do not share variables
- The processes access common resources, but in a critical section
- Assume that there is only one critical section (can extend to more than one critical section)
- Assumption:
 - system is asynchronous
 - processes do not fail and message delivery is reliable
 - any message sent is eventually delivered intact, exactly once
- The application-level protocol for executing a critical section is as follows:
enter() // enter critical section – block if necessary
resourceAccesses() // access shared resources in critical section
exit() // leave critical section – other processes may now enter

Essential requirements for mutual exclusion are as follows:

ME1: (safety)

At most one process may execute in the critical section (CS) at a time

ME2: (liveness)

Requests to enter and exit the critical section eventually succeed

- **Condition ME2 implies freedom from both deadlock and starvation**
- Even without a deadlock, a poor algorithm might lead to *starvation: the indefinite* postponement of entry for a process that has requested it.
- The absence of starvation is a ***fairness condition***.

- Another fairness issue: order in which processes enter the critical section
- Cannot **order entry** to the critical section by time due to **absence of global clocks**
- **Happened-before ordering** may be used

ME3: (-> ordering) If one request to enter the CS happened-before another, then entry to the CS is granted in that order

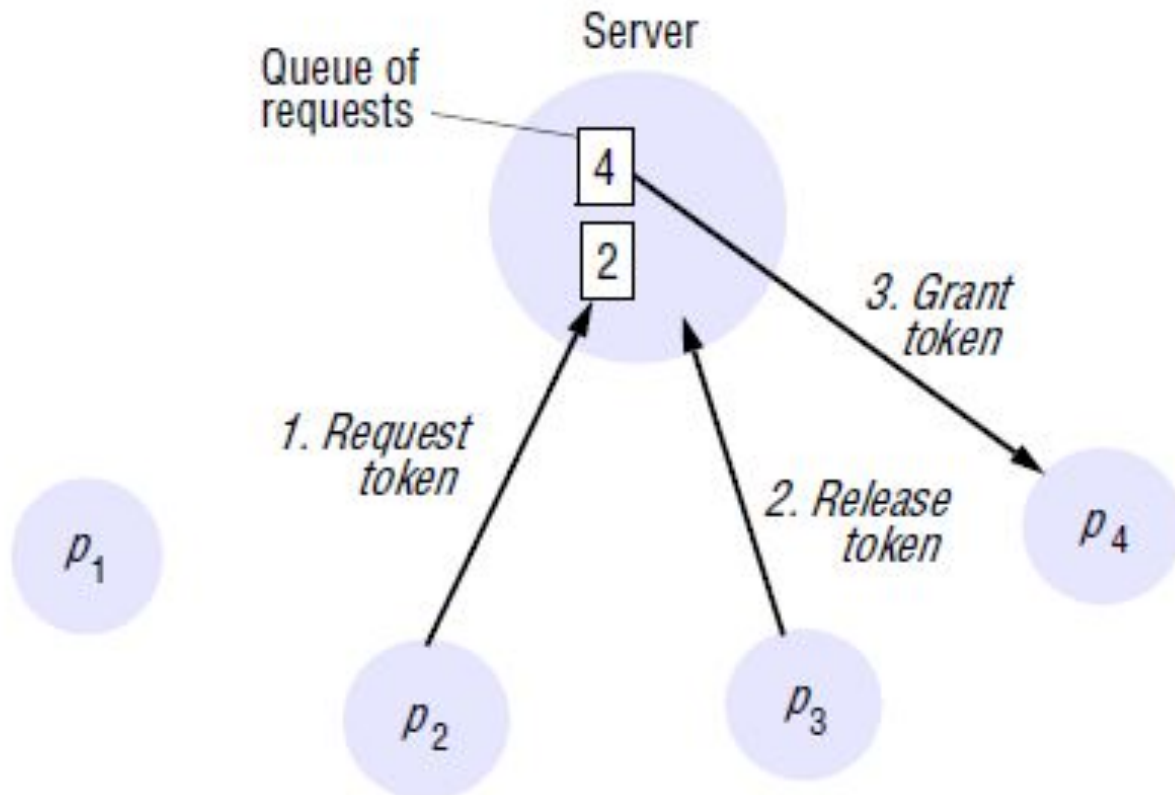
- Hence it is not possible for a process to enter the critical section more than once while another waits to enter

- We evaluate the performance of algorithms for mutual exclusion according to the following criteria:
 - ✓ the *bandwidth consumed*
(number of messages sent in each entry and exit operation)
 - ✓ the client *delay incurred* by a process at each entry and exit operation
 - ✓ the algorithm's effect upon the *throughput of the system*
Rate at which the collection of processes as a whole can access the critical section
- The effect is measured using the *synchronization delay between one process exiting the critical section and the next process entering it*
- Throughput is greater when the synchronization delay is shorter

The central server algorithm

- Simplest way to achieve mutual exclusion
- Server grants permission to enter the critical section

Server managing a mutual exclusion token for a set of processes




- To enter a critical section, a process sends a request message to the server and awaits a reply from it.
- Conceptually, the reply constitutes a token signifying permission to enter the critical section.
- If no other process has the token at the time of the request, then the server replies immediately, granting the token.
- If the token is currently held by another process, then the server does not reply, but queues the request.
- When a process exits the critical section, it sends a message to the server, giving it back the token.
- If the queue of waiting processes is not empty, then the server chooses the oldest entry in the queue, removes it and replies to the corresponding process.
- The chosen process then holds the token.
- Safety and liveness conditions are met by this algorithm.

Performance of this algorithm:

- Entering the critical section takes two messages (a *request* followed by a *grant*)
- *Delays the requesting process by the time required for this round-trip*
- Exiting the critical section takes one *release message*
- This does not delay the exiting process
- The server may become a performance bottleneck for the system
- The synchronization delay is the time taken for a round-trip:
 - *a release message to the server*
 - *a grant message to the next process to enter the critical section*

A ring-based algorithm

- Simple way to arrange mutual exclusion between the N processes
- *No additional process required to arrange them in a logical ring*
- Requirement: *a communication channel from each process P_i has to the next process in the ring, $P(i + 1) \bmod N$*
- *Exclusion is conferred by obtaining a token in the form of a message passed from process to process in a single direction*
- The ring topology may be unrelated to the physical interconnections between the underlying computers




- If a process does not require to enter the critical section when it receives the token, then it immediately forwards the token to its neighbour.
- A process that requires the token waits until it receives it, but retains it.
- To exit the critical section, the process sends the token on to its neighbour.
- Conditions **ME1 and ME2 are met** by this algorithm
- But, the token is **not necessarily obtained in happened-before order**
- Processes may exchange messages independently of the rotation of the token

- This algorithm continuously consumes network bandwidth (except when a process is inside the critical section)
- Delay experienced by a process requesting entry to the critical section: between 0 messages (just received the token) and *N messages (just passed on the token)*
- *To exit the* critical section requires only one message
- Synchronization delay between one process's exit from the critical section and the next process's entry = From 1 to *N message transmissions*

An algorithm using multicast and logical clocks (Ricart and Agrawala's)

Multicast synchronization



Ricart and Agrawala's Algorithm

On initialization

state := RELEASED;

To enter the section

state := WANTED;

Multicast *request* to all processes;

T := request's timestamp;

Wait until (number of replies received = (*N* − 1));

state := HELD;

Request processing deferred here

On receipt of a request $\langle T_i, p_i \rangle$ *at* p_j ($i \neq j$)

if (*state* = HELD or (*state* = WANTED and $(T, p_j) < (T_i, p_i)$))

then

 queue *request* from p_i without replying;

else

 reply immediately to p_i ;

end if

To exit the critical section

state := RELEASED;

reply to any queued requests;

Maekawa's voting algorithm

- Not necessary for all of its peers to grant it access
- Need only obtain permission to enter from *subsets of their peers*
- *The subsets* used by any two processes should overlap
- Processes vote for one another to enter the critical section
- A 'candidate' process must collect sufficient votes to enter
- Processes in the intersection of two sets of voters ensure the safety property ME1

- Let there be 7 processes- 0,1,2,3,4,5,6

S_0	=	{0, 1, 2}
S_1	=	{1, 3, 5}
S_2	=	{2, 4, 5}
S_3	=	{0, 3, 4}
S_4	=	{1, 4, 6}
S_5	=	{0, 5, 6}
S_6	=	{2, 3, 6}

- Voting set V_i with each process p_i ($i = 1, 2, \dots, N$), where V_i is a subset of $\{p_1, p_2, \dots, p_N\}$
- The sets V_i are chosen so that, for all $i, j = 1, 2, \dots, N$:
 - $p_i \in V_i$
 - $V_i \cap V_j \neq \emptyset$ – there is at least one common member of any two voting sets
 - $|V_i| = K$ – to be fair, each process has a voting set of the same size
 - Each process p_j is contained in M of the voting sets V_i
- Optimal solution:
 - minimizes K and allows the processes to achieve mutual exclusion
 - $K \sim \sqrt{N}$ and $M = K$ (so that each process is in as many of the voting sets as there are elements in each one of those sets).
- It is nontrivial to calculate the optimal sets R_i

On initialization

state := RELEASED;

voted := FALSE;

For p_i to enter the critical section

state := WANTED;

Multicast *request* to all processes in V_i ;

Wait until (number of replies received = K);

state := HELD;

On receipt of a request from p_i at p_j

if (*state* = HELD or *voted* = TRUE)

then

queue *request* from p_i without replying;

else

send *reply* to p_i ;

voted := TRUE;

end if

For p_i to exit the critical section

state := RELEASED;

Multicast *release* to all processes in V_i ;

On receipt of a release from p_i at p_j

if (queue of requests is non-empty)

then

remove head of queue – from p_k , say;

send *reply* to p_k ;

voted := TRUE;

else

voted := FALSE;

end if

- To obtain entry to the critical section, a process p_i sends request messages to all K members of V_i (including itself).
- p_i cannot enter the critical section until it has received all K reply messages.
- When a process p_j in V_i receives p_i 's request message, it sends a reply message immediately, unless either its state is HELD or it has already replied ('voted') since it last received a release message.
- Otherwise, it queues the request message (in the order of its arrival) but does not yet reply.
- When a process receives a release message, it removes the head of its queue of outstanding requests (if the queue is nonempty) and sends a reply message (a 'vote') in response to it.
- To leave the critical section, p_i sends release messages to all K members of V_i (including itself).

- This algorithm achieves the safety property, ME1

Unfortunately, the algorithm is deadlock-prone:

- Consider three processes, $p1$, $p2$ and $p3$, with
 - $V1 = \{ p1, p2 \}$
 - $V2 = \{ p2, p3 \}$
 - $V3 = \{ p3, p1 \}$
- *If the three processes concurrently request entry to the critical section*
- *Each process has received one out of two replies, and none can proceed*

- The algorithm can be adapted so that it becomes deadlock-free
- Processes may queue outstanding requests in happened-before order
- ME3 is also satisfied
- The algorithm's bandwidth utilization is
 - $2N$ messages per entry to the critical section
 - N messages per exit (assuming no hardware multicast facilities)

ELECTIONS

- Algorithm for choosing a unique process to play a particular role is called an *election algorithm*
 - For example, in a variant of central-server algorithm for mutual exclusion, the 'server' is chosen from among the processes $p_i, (i = 1, 2, \dots, N)$ that need to use the critical section
- An election algorithm is needed for this choice
- It is essential that all the processes agree on the choice
- Afterwards, if the server wishes to retire then another election is required to choose a replacement

- Process **calls the election**- takes an action that initiates a run of the election algorithm.
- An individual process does not call more than one election at a time
- **N processes could call N concurrent elections**
- At any point in time, a process p_i is either a **participant** or a **non-participant**
- The **elected process** must be **unique**
 - For example, two processes could decide independently that a coordinator process has failed, and both call elections
- The elected process be chosen as **the one with the largest identifier**
- The 'identifier' may be any useful value, unique and totally ordered


- Each process *has a variable elected_i*, which will contain the identifier of the elected process
- Variable elected_i set to the special value '⊥' initially to denote that it is not yet defined.

Requirements during any particular run of the algorithm:

- **E1: (safety)** A participant process *p_i* has *elected_i = ⊥* or *elected_i = P*, where *P* is non-crashed process chosen at the end of the run with the largest identifier
- **E2: (liveness)** All processes *p_i* participate and eventually either set *elected_i ≠ ⊥* – or crash
- There may be processes *p_j* that are not yet participants, which record in *elected_j* the identifier of the previous elected process

A RING-BASED ELECTION ALGORITHM

- Each process p_i has a communication channel to the next $(p_i + 1 \bmod N)$
- Assumption: no failures occur, and the system is asynchronous
- Goal: To elect a single process called the *coordinator*, which is the process with the largest identifier



- Initially, every process is marked as a *non-participant*
- *Any process* can begin an election
 - marks itself as a *participant*
 - *placing its identifier* in an *election message*
 - *sending it to its clockwise neighbour*
- When a process receives an *election message*
 - *it compares the identifier in the message* with its own
 - If the arrived identifier is greater- forwards the message to its neighbour
 - If the arrived identifier is smaller and the receiver is not a *participant*- *then* it substitutes its own identifier in the message and forwards it
 - Does not forward the message if it is already a *participant*
- *On forwarding an election message* the process marks itself as a *participant*

- If the received identifier is that of the receiver itself
 - Its identifier must be the greatest, and it becomes the coordinator
- The coordinator
 - marks itself as a *non-participant once more*
 - *sends an elected message to its neighbour*, announcing its election and enclosing its identity
- When a process *pi* receives an elected message:
 - *it marks itself as a nonparticipant*
 - sets its variable *electedi* to the identifier in the message
 - *unless it is the new coordinator*, forwards the message to its neighbour

Condition E1 is met:

- All identifiers are compared, since a process must receive its own identifier back before sending an *elected message*.
- *For any* two processes the one with the larger identifier will not pass on the other's identifier
- It is therefore impossible that both should receive their own identifier back.

Condition E2 is met:

- Follows immediately from the guaranteed traversals of the ring (if there are no failures)
- The *non-participant and participant states*
 - *Eliminate duplicate messages* as soon as possible, and always before the 'winning' election result has been announced

Performance if only a single process starts an election:

- Worst-performing case - its anti-clockwise neighbour has the highest identifier
 - A total of $N - 1$ *messages* are then required to reach this neighbour
 - Its election announced completed with another N *messages*
 - *The elected* message is then sent N *times*
 - *Total* = $3N - 1$ *messages in all*
 - *The turnaround time* is also $3N - 1$
- Ring-based algorithm is useful for understanding the properties of election algorithms
- Not practical as it tolerates no failures
- A reliable failure detector may help to reconstitute the ring when a process crashes

THE BULLY ALGORITHM

- The bully algorithm allows processes to crash during an election
- It assumes that message delivery between processes is reliable
- Assumes that the **system is synchronous**:
it uses timeouts to detect a process failure
- Another difference
 - The ring-based algorithm assumed that processes have minimal *a priori knowledge of one* another, communicates with only its neighbour
 - The bully algorithm assumes that **each process knows which processes have higher identifiers**, and can communicate with all such processes

There are three types of message in this algorithm:

- **election message** is sent to announce an election
 - **answer message** is sent in response to an election message
 - **coordinator message** is sent to announce the identity of the elected process – the new ‘coordinator’.
-
- A process begins an election when it notices, through timeouts, that the coordinator has failed
 - Several processes may discover this concurrently

- Since the system is synchronous, we can construct a reliable failure detector
- Maximum message transmission delay = T_{trans}
- Maximum delay for processing a message $T_{process}$.
- **Round trip time $T = 2T_{trans} + T_{process}$**
 - Upper bound on the time that can elapse between sending a message to another process and receiving a response.
- If no response arrives within time T -
 - local **failure detector** can report that intended recipient has failed


- The process that **knows it has the highest identifier** can elect itself as the coordinator simply by sending a *coordinator message to all processes with lower identifiers*.
- On the other hand, a process with a **lower identifier** can **begin an election** by sending an *election message to those processes that have a higher identifier* and awaiting *answer messages* in response.
- *If none arrives within time T* , the process *considers itself* the coordinator and sends a *coordinator message to all processes with lower identifiers* announcing this.
- **Otherwise**, the process *waits a further period T' for a coordinator message* to arrive from the new coordinator.
- *If none arrives*, it begins another election.

- If a process p_i *receives a coordinator message*, it sets its variable *elected* to the identifier of the coordinator contained within it and treats that process as the coordinator.
- If a process receives an *election message*, it sends back an *answer message* and begins another election – unless it has begun one already.

When a process is started, to replace a crashed process, it begins an election

- If it has the **highest process identifier**, then it will **decide that it is the coordinator** and **announce this to the other processes**
- Thus it will become the coordinator, **even though the current coordinator is functioning**
- It is for this reason that the algorithm is called the '**bully**' algorithm

The bully algorithm



The election of coordinator p_2 , after the failure of p_4 and then p_3

There are four processes, $p_1 - p_4$.

- **Stage1**: Process p_1 detects the failure of the coordinator p_4 and announces an election.

On receiving an election message from p_1 , processes p_2 and p_3 send answer messages to p_1 and begin their own elections;

- **Stage 2**: p_3 sends an answer message to p_2 , but p_3 receives no answer message from the failed process p_4 .

It therefore decides that it is the coordinator.

- **Stage 3**: But before it can send out the coordinator message, it too fails. When p_1 's timeout period T' expires (which we assume occurs before p_2 's timeout expires), it deduces the absence of a coordinator message and begins another election.
- **Stage 4**: Eventually, p_2 is elected coordinator.

- **Liveness condition E2 is met** by assuming reliable message delivery
- Condition E1 is met if **no process is replaced**
- *Safety condition E1 not met if*
 - *processes that have crashed are replaced by processes with the same identifiers*
 - the assumed timeout values turn out to be inaccurate (suppose p3 was just slow)

Performance of the algorithm:

- Best case - the process with the second-highest identifier notices the coordinator's failure
 - Then it can immediately elect itself and send $N - 2$ *coordinator messages*.
 - *The turnaround time is one message.*
- Worst case - process with lowest identifier first detects the coordinator's failure
 - The bully algorithm requires $O(N^2)$ *messages*
 - For then $N - 1$ processes altogether begin elections, each sending messages to processes with higher identifiers.

- This algorithm clearly meets the **liveness condition E2**, by the assumption of reliable message delivery.
- And if **no process is replaced**, then the algorithm **meets condition E1**.
- *It is impossible for two processes to decide that they are the coordinator*, since the process with the lower identifier will discover that the other exists and defer to it.
- But the algorithm is ***not guaranteed to meet the safety condition E1 if processes that have crashed are replaced by processes with the same identifiers.***
- A process that replaces a crashed process *p* may decide that it has the highest identifier just as another process (which has detected *p*'s crash) decides that it has the highest identifier.
- **Two processes will therefore announce themselves as the coordinator concurrently.**
- Unfortunately, there are no guarantees on message delivery order, and the recipients of these messages may reach different conclusions on which is the coordinator process.
- Furthermore, condition E1 may be broken if the assumed timeout values turn out to be inaccurate – that is, if the processes' failure detector is unreliable.

- Suppose that either *p3 had not failed but was just running unusually slowly* (that is, that the assumption that the system is synchronous is incorrect), or that *p3 had failed but was then replaced*.
- *Just as p2 sends its coordinator message, p3 (or its replacement) does the same.*
- *p2 receives p3 's coordinator message after it has sent its own and so sets $elected2 = p3$.*
- *Due to variable message transmission delays, p1 receives p2 's coordinator message after p3 's and so eventually sets $elected1 = p2$.*
- *Condition E1 has been broken.*

THANK YOU!